

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/5213438>

How to Forecast Long-Run Volatility: Regime Switching and the Estimation of Multifractal Processes

Article in *Journal of Financial Econometrics* · February 2004

DOI: 10.1093/jjfinec/nbh003 · Source: RePEc

CITATIONS

156

READS

186

1 author:



[Laurent E. Calvet](#)

Edhec Business school

78 PUBLICATIONS 2,410 CITATIONS

SEE PROFILE

Some of the authors of this publication are also working on these related projects:



Robust filtering [View project](#)



Asset pricing with multifrequency states [View project](#)

How to Forecast Long-Run Volatility: Regime Switching and the Estimation of Multifractal Processes

LAURENT E. CALVET

Harvard University and NBER

ADLAI J. FISHER

University of British Columbia

ABSTRACT

We propose a discrete-time stochastic volatility model in which regime switching serves three purposes. First, changes in regimes capture low-frequency variations. Second, they specify intermediate-frequency dynamics usually assigned to smooth autoregressive transitions. Finally, high-frequency switches generate substantial outliers. Thus a single mechanism captures three features that are typically viewed as distinct in the literature. Maximum-likelihood estimation is developed and performs well in finite samples. Using exchange rates, we estimate a version of the process with four parameters and more than a thousand states. The multifractal outperforms GARCH, MS-GARCH, and FIGARCH in- and out-of-sample. Considerable gains in forecasting accuracy are obtained at horizons of 10 to 50 days.

KEYWORDS: forecasting, long memory, Markov-switching multifractal (MSM), closed-form likelihood, scaling, stochastic volatility, volatility component, Vuong test

Over the past 15 years, stochastic regime switching (Hamilton 1989, 1990) has proven to be extremely useful for modeling economic and financial time series.¹

We received useful comments from John Campbell, René Garcia (the editor), Guido Kuersteiner, Nour Meddahi, and two anonymous referees. We also thank T. Andersen, T. Bollerslev, F. Diebold, R. Engle, J. Hamilton, J. MacKinnon, N. Shephard, D. Smith, R. Tsay, and seminar participants at the University of British Columbia, University of Chicago, University of Pennsylvania, 2002 NSF-NBER Time Series Conference at the Wharton School, 2002 CIRANO Extreme Events in Finance Conference in Montréal, and 2003 North American Summer Meeting of the Econometric Society at Northwestern University. This paper is a revised version of an earlier manuscript, "Regime-Switching and the Estimation of Multifractal Processes" (Calvet and Fisher 2002). Excellent research assistance was provided by Xifeng Diao. We are very appreciative of financial support provided for this project by the Social Sciences and Humanities Research Council of Canada under grant 410-2002-0641. Address correspondence to Department of Economics, Cambridge, MA 02138, or e-mail: lcalvet@aya.yale.edu; or Finance Division, Sauder School of Business, 2053 Main Mall, Vancouver, BC, Canada V6T 1Z2, or e-mail: adlai.fisher@sauder.ubc.ca.

¹ The likelihood-based estimation of Markov switching processes was developed by Lindgren (1978) and Baum et al. (1980) in the statistics literature. A seminal series of articles by Hamilton (1989, 1990) introduced

While the theoretical formulation is very general, empirical researchers most commonly apply this approach to low-frequency variations and rely on other techniques for shorter-run dynamics. For example, Markov switching autoregressive conditionally heteroscedastic (ARCH) and generalized autoregressive conditionally heteroscedastic (GARCH) processes separately specify regime shifts at low frequencies, smooth autoregressive volatility transitions at midrange frequencies, and a thick-tailed conditional distribution of returns at high frequency [Cai (1994), Hamilton and Susmel (1994), Gray (1996), Klaassen (2002)]. In this article, we propose an alternative volatility model based on pure regime switching at all frequencies.

Previous empirical applications typically employ only a small number of discrete states. This partly stems from the common view that regime switches occur infrequently. In a general formulation, a more practical limitation is that the number of parameters grows quadratically with the cardinality of the state space. Restrictions on switching probabilities offer a natural solution, as pursued for example by Bollen, Gray, and Whaley (2000) in a four-regime model.² We extend this approach by considering a tight set of restrictions inspired by the multifractal literature. This permits the routine estimation of good-performing models with more than a thousand states, a dense transition matrix, and only four parameters.

Our specification is particularly influenced by an innovation in multifractal modeling from Calvet and Fisher (2001). This earlier theoretical research uses Markov switching to develop the first time-stationary formulation of multifractal diffusions and also provides a weakly convergent sequence of discrete filters. We now propose a variant of these filters to directly model financial series in discrete time. In this framework, total volatility is the multiplicative product of a large but finite number of random components. We assume for simplicity that these components are first-order Markov and identical except for time scale. The components have identical marginal distributions and differ only in their switching probabilities. The specification is completed by assuming that the progression of switching probabilities is approximately geometric. The Markov construction delivers a parsimonious stochastic volatility model with a closed-form likelihood. The multifractal structure also generates substantial outliers, long-memory features in volatility, and a decomposition of volatility into components with

these processes to econometrics and spurred the development of a large body of research. Contributions to the original version of the model advance estimation and testing [Hansen (1992), Shephard (1994), Albert and Chib (1993), Garcia (1998)] and investigate a wide range of empirical applications [e.g., Hamilton (1988), Garcia and Perron (1996)]. The approach has been extended to incorporate GARCH transitions [Cai (1994), Hamilton and Susmel (1994), Kim (1994), Gray (1996), Kim and Nelson (1999), Klaassen (2002)], vector processes [Hamilton and Lin (1996), Hamilton and Perez-Quiros (1996)], and time-varying transition probabilities [Diebold, Lee, and Weinbach (1994), Durland and McCurdy (1994), Filardo (1994), Perez-Quiros and Timmerman (2000)]. See Hamilton and Raj (2002) for a recent survey.

² Duration-dependent Markov switching models also use restrictions on state parameters and switching probabilities. For example, Maheu and McCurdy (2000) expand a two-state model by conditioning the volatility level and switching probability on duration in the state. The resulting transition matrix is sparse, and the system either progresses to the next duration of the same state or the first duration of the other state.

heterogenous decay rates. We accordingly call our model the Markov switching multifractal (MSM).

An empirical investigation of four daily exchange rate series shows that MSM performs well in comparison with the Student-GARCH(1,1) of Bollerslev (1987), the Markov switching GARCH (MS-GARCH) of Klaassen (2002), and the fractionally integrated GARCH (FIGARCH) of Baillie, Bollerslev, and Mikkelsen (1996). The choice of alternative processes is guided by several considerations. First, the multifractal easily permits maximum-likelihood (ML) estimation and analytical multistep forecasting. We therefore compare it to models with the same appealing features. GARCH(1,1) is an obvious choice because it is a leading model for volatility forecasting [see, e.g., Akgiray (1989), Pagan and Schwert (1990), West and Cho (1995), Andersen and Bollerslev (1998), and Hansen and Lunde (2001)]. MS-GARCH combines regime shifts with smooth weighting of past shocks and thus represents a compromise between the two approaches.³ Like the multifractal, FIGARCH permits long memory in volatility and is thus consistent with empirical evidence on exchange rates [Baillie, Bollerslev, and Mikkelsen (1996)].

The multifractal process compares favorably with GARCH(1,1) both in- and out-of-sample. Our model has a higher likelihood in-sample for all currencies, and the statistical significance of these differences is confirmed by a Heteroskedastic and Autocorrelation Consistent (HAC) version of the Vuong (1989) test. Since both models have the same number of parameters, the multifractal is also preferred by standard selection criteria. Analogous results are obtained out-of-sample. While one-day forecasts from the two models perform similarly, the multifractal dominates at longer horizons. The difference is most pronounced for 20 and 50 days. For example, in the case of the British pound, the 50-day forecasting R^2 is 27.3% for MSM as compared to -2.6% for GARCH(1,1). Similar gains are obtained for other currencies. The empirical evidence thus shows that the multifractal model improves on GARCH(1,1) both in- and out-of-sample.

The MS-GARCH process gives substantially better fit than GARCH(1,1) in-sample. This is partly attributable to a larger number of parameters, and suggests the possibility of overfitting. Using a Bayesian information criterion (BIC), the multifractal is statistically indistinguishable from MS-GARCH in-sample for all four currencies. Out-of-sample, MS-GARCH is also comparable to the multifractal at short horizons, but substantially dominated at longer horizons. Thus MS-GARCH is overall dominated by MSM.

The multifractal also compares favorably with FIGARCH in- and out-of-sample. While MSM and FIGARCH have the same number of parameters, the estimated likelihoods are substantially higher with our model for all currencies. Out-of-sample, FIGARCH tends to provide poorer forecasts than the multifractal, even at long horizons. For instance, with the Japanese yen and British pound, the 50-day forecasting R^2 is negative for FIGARCH compared with values larger than 20% for our model. The multifractal is overall the best-performing process.

³ The original MS-GARCH process of Gray (1996) does not conveniently permit multistep forecasting, and we therefore consider the variant formulation introduced for this purpose by Klaassen (2002).

A pure Markov switching model thus captures the same dynamics that in previous literature have required not only regime switching but also linear GARCH transitions and a thick-tailed conditional distribution of returns. It is striking that a single mechanism can play all three of these roles so effectively. Our innovation that achieves this surprising economy of modeling technique is based on scale invariance. We make this principle operational by introducing a pure Markov switching formulation where scale invariance helps to specify the parameters and transitions of a high-dimensional state space.

Our article contributes to the multifractal literature by offering a convenient time-series construction accompanied by effective estimation and testing methods. Calvet, Fisher, and Mandelbrot (1997) introduce the Multifractal Model of Asset Returns (MMAR), a class of diffusions that capture the outliers, moment scaling, and long memory in volatility exhibited by many financial time series. While providing an excellent fit to many aspects of financial data, the MMAR uses a combinatorial construction that is not particularly well suited to econometrics. In particular, regime changes take place at predetermined dates, making the model non stationary. Early estimation efforts have thus focused on unconditional moments of returns.⁴ Calvet and Fisher (2002a) use moment scaling to develop an estimator and diagnostic tests of the model. This previous empirical work is constrained in many ways by the nonstationarity of the MMAR. We now employ a stationary Markov switching formulation and develop ML estimation, which is new to the literature on multifractal measures and processes.

Section 1 presents the discrete-time model. Section 2 develops the ML estimator and assesses its accuracy in Monte Carlo simulations. Section 3 discusses estimation results for four exchange rates. Section 4 compares our model with alternative processes both in- and out-of-sample. Section 5 concludes. All proofs are given in the appendix.

1 THE MARKOV SWITCHING MULTIFRACTAL

This section develops MSM, a discrete-time Markov process with multifrequency stochastic volatility. The process has a finite number \bar{k} of latent volatility state variables, each of which corresponds to a different frequency.

1.1 Stochastic Volatility

We consider an economic series X_t defined in discrete time on the regular grid $t = 0, 1, 2, \dots, \infty$. In applications, X_t will be the log-price of a financial asset or exchange rate. Define the innovations $x_t \equiv X_t - X_{t-1}$. A common modeling methodology assumes that the system is hit every period by a single shock that

⁴ The MMAR implies that return moments vary as power functions of the observation interval, which is consistent with many financial series. Calvet, Fisher, and Mandelbrot (1997) and Calvet and Fisher (2002a) find evidence of moment scaling in powers of the absolute value of returns. Further evidence is provided by Andersen et al. (2001) and Barndorff-Nielsen and Shephard (2003). See LeBaron (2001) for a discussion of robustness.

progressively phases out over time [e.g., Engle (1982)]. We consider instead an economy with \bar{k} components $M_{1,t}, M_{2,t}, \dots, M_{\bar{k},t}$, which decay at heterogeneous frequencies $\gamma_1, \dots, \gamma_{\bar{k}}$. Such a model could be very unwieldy as the number \bar{k} becomes very large. We will see, however, that the process can be parsimoniously specified by a small set of parameters.

We model the returns $x_t \equiv X_t - X_{t-1}$ by

$$x_t = \sigma(M_{1,t}M_{2,t} \dots M_{\bar{k},t})^{1/2} \varepsilon_t, \quad (1)$$

where σ is a positive constant and the random variables $\{\varepsilon_t\}$ are i.i.d. standard Gaussians $\mathcal{N}(0, 1)$. The random *multipliers* or *volatility components* $M_{k,t}$ are persistent, nonnegative, and satisfy $\mathbb{E}(M_{k,t}) = 1$. We consider for simplicity that the multipliers $M_{1,t}, M_{2,t}, \dots, M_{\bar{k},t}$ at a given time t are statistically independent. The parameter σ is then equal to the unconditional standard deviation of the innovation x_t .

Equation (1) defines a stochastic volatility model $x_t = \sigma_t \varepsilon_t$ with the multiplicative structure $\sigma_t = \sigma(M_{1,t}M_{2,t} \dots M_{\bar{k},t})^{1/2}$. We conveniently stack the period t volatility components into a vector

$$M_t = (M_{1,t}, M_{2,t}, \dots, M_{\bar{k},t}).$$

For any $m = (m_1, \dots, m_{\bar{k}}) \in \mathbb{R}^{\bar{k}}$, let $g(m)$ denote the product $\prod_{i=1}^{\bar{k}} m_i$. Volatility at time t is then $\sigma_t = \sigma[g(M_t)]^{1/2}$.

The properties of volatility are driven by the stochastic dynamics of the vector M_t . We assume for parsimony that M_t is first-order Markov. This facilitates the simulation of $\{x_t\}$ and permits ML estimation. It is natural to call M_t the *volatility state vector* and each component $M_{k,t}$ a *state variable*. The econometrician observes the returns $x_t = \sigma[g(M_t)]^{1/2} \varepsilon_t$, but not the vector M_t itself. The vector M_t is therefore latent, and must be inferred recursively by Bayesian updating.

Each $M_{k,t}$ follows a process that is identical except for time scale. Assume that the volatility state vector has been constructed up to date $t - 1$. For each $k \in \{1, \dots, \bar{k}\}$, the next period multiplier $M_{k,t}$ is drawn from a fixed distribution M with probability γ_k , and is otherwise equal to its previous value: $M_{k,t} = M_{k,t-1}$. The dynamics of $M_{k,t}$ can be summarized as

$$\begin{array}{ll} M_{k,t} \text{ drawn from distribution } M & \text{with probability } \gamma_k \\ M_{k,t} = M_{k,t-1} & \text{with probability } 1 - \gamma_k. \end{array}$$

The switching events and new draws from M are assumed to be independent across k and t . The volatility components $M_{k,t}$ thus differ in their transition probabilities γ_k , but not in their marginal distribution M . These features greatly contribute to the parsimony of the model.

The transition probabilities $\gamma \equiv (\gamma_1, \gamma_2, \dots, \gamma_{\bar{k}})$ are specified as

$$\gamma_k = 1 - (1 - \gamma_1)^{(b^{k-1})}, \quad (2)$$

where $\gamma_1 \in (0, 1)$ and $b \in (1, \infty)$. This specification is introduced in Calvet and Fisher (2001) in connection with the discretization of a Poisson arrival process.

Since $1 - \gamma_k = (1 - \gamma_1)^{(b^k - 1)}$, the logarithms of staying probabilities are exponentially decreasing with k . Consider a process with very persistent components and thus a very small parameter γ_1 . For small values of k , the quantity $\gamma_1 b^{k-1}$ remains small and the transition probability satisfies

$$\gamma_k \sim \gamma_1 b^{k-1}.$$

The transition probabilities of low-frequency components grow approximately at geometric rate b . At higher frequencies ($\gamma_k \sim 1$), the rate of increase slows down and Equation (2) guarantees that the parameter γ_k remains less than one. In empirical applications, it is numerically convenient to estimate parameters of the same magnitude. Since $\gamma_1 < \dots < \gamma_{\bar{k}} < 1 < b$, we choose $(\gamma_{\bar{k}}, b)$ to specify the set of transition probabilities.

The integer \bar{k} determines the number of volatility frequencies and its choice is viewed as a model selection problem. The multifractal construction imposes only minimal restrictions on the marginal distribution of the multipliers: $M \geq 0$ and $\mathbb{E}(M) = 1$. This allows flexible parametric or even nonparametric specifications of M . For simplicity, however, this article focuses on the simple case where M is a binomial random variable taking values m_0 or $2 - m_0$ with equal probability. The full parameter vector is then

$$\psi \equiv (m_0, \sigma, b, \gamma_{\bar{k}}) \in \mathbb{R}_{+}^4,$$

where m_0 characterizes the distribution of the multipliers, σ is the unconditional standard deviation of returns, and b and $\gamma_{\bar{k}}$ define the set of switching probabilities. In Section 2, we will develop and empirically implement the ML estimation of this vector.

We call this process the Markov Switching Multifractal. The notation $\text{MSM}(\bar{k})$ refers to versions of the model with \bar{k} frequencies. Economic intuition and earlier work suggest that the multiplicative structure of Equation (1) is appealing to model the high-variability and high-volatility persistence exhibited by financial time series. When a low-level multiplier changes, volatility varies discontinuously and has strong persistence. In addition, high-frequency multipliers produce substantial outliers.

1.2 Properties

The MSM construction permits low-frequency regime shifts, and thus long volatility cycles in sample paths. We will see that in exchange rate series, the duration of the most persistent component, $1/\gamma_1$, is typically of the same order as the length of the data. Estimated processes thus generate volatility cycles with periods proportional to the sample size, a property also apparent in the sample paths of long-memory processes.

Long memory is often defined by a hyperbolic decline in the autocovariance function as the lag goes to infinity. Fractionally integrated processes generate such patterns by assuming that an innovation linearly affects future periods at a hyperbolically declining weight. As a result, fractional integration tends to

produce smooth volatility processes. By contrast, our approach generates long cycles with a switching mechanism that also gives abrupt volatility changes. The combination of long-memory behavior with sudden volatility movements has a natural appeal for financial econometrics.

In earlier work, we proposed a definition of long memory that applies to continuous-time processes defined on a bounded time domain. This definition is based on increments over progressively smaller intervals. Our discrete-time process can also generate a hyperbolic decline in autocovariances for a range of lags. For every moment $q \geq 0$ and every integer n , let $\rho_q(n) = \text{corr}(|x_t|^q, |x_{t+n}|^q)$ denote the autocorrelation in levels. We choose a fixed vector ψ and consider the positive parameter $\delta(q) = \log_b(\mathbb{E}(M^q)/[\mathbb{E}(M^{q/2})]^2)$. Consider two arbitrary numbers α_1 and α_2 in the open interval $(0, 1)$. The set of integers $I_{\bar{k}} = \{n : \alpha_1 \log_b(b^{\bar{k}}) \leq \log_b n \leq \alpha_2 \log_b(b^{\bar{k}})\}$ contains a large range of intermediate lags. We show this in the appendix

Proposition 1. *The autocorrelation in levels satisfies*

$$\sup_{n \in I_{\bar{k}}} \left| \frac{\ln \rho_q(n)}{\ln n^{-\delta(q)}} - 1 \right| \rightarrow 0$$

as $\bar{k} \rightarrow +\infty$.

MSM thus mimics the hyperbolic autocovariograms $\ln \rho_q(n) \sim -\delta(q) \ln n$ exhibited by many financial series [e.g., Dacorogna et al. (1993), Ding, Granger, and Engle (1993), Baillie, Bollerslev, and Mikkelsen (1996), Gouriéroux and Jasiak (2002)]. This result complements earlier research that has emphasized the difficulty of distinguishing between long memory and structural change in finite samples [e.g., Bhattacharya, Gupta, and Weymire (1983), Hidalgo and Robinson (1996), Diebold and Inoue (2001)]. MSM illustrates that a Markov chain regime-switching model can exhibit one of the defining features of long memory, a hyperbolic decline of the autocovariogram.

A representative return series is illustrated in Figure 1. The graph reveals large heterogeneity in volatility levels and substantial outliers. This is notable




Figure 1 Simulated multifractal process. This figure shows simulated log-price differences of a multifractal process. The process has $\bar{k} = 8$ frequencies and parameter values $m_0 = 1.4$, $\sigma = 0.5$, $\gamma_{\bar{k}} = 0.95$, and $b = 3$. These parameter values are roughly consistent with estimates that are found to provide a good description of several exchange rate series in later sections of the article.

since the return process has by construction finite moments of every order. In Calvet and Fisher (2001), we introduced Paretian tails by considering i.i.d. shocks of the form $\varepsilon_t = Z_t \Omega_t$, where $\{Z_t\}$ are i.i.d. standard Gaussians and $\{\Omega_t\}$ are i.i.d. draws from a distribution Ω with Paretian tails. The random variable Ω is the limit distribution of renormalized quadratic variation over a fixed horizon, and is fully specified by m_0 , b , and $\gamma_{\bar{k}}$. The specification $\varepsilon_t = Z_t \Omega_t$ represents an interesting direction for future research, but this article focuses on the Gaussian case for several reasons. First, the likelihood is then available in closed form. Second, we will show that even when ε_t is Gaussian, high-frequency regime switches are sufficient to mimic in finite samples the heavy tails exhibited by the data. Finally, the basic specification of our process performs well relative to existing competitors and provides a useful benchmark for future refinements.

Another interesting property of MSM is that when $\bar{k} \rightarrow \infty$, the limiting continuous-time process lies outside the class of Itô diffusions. The sample paths are continuous but exhibit a high degree of heterogeneity in local behavior, which is characterized by a continuum of local Hölder exponents in any finite time interval. We refer the reader to Calvet and Fisher (2001, 2002a) for a full development of the continuous-time multifractal limit.

2 ML ESTIMATION

When the multiplier M has a discrete distribution, there exist a finite number of volatility states. Standard filtering methods then provide the likelihood function in closed form.

2.1 Updating the State Vector

We assume in the rest of this article that the multiplier M has a finite number b_m of values. The volatility state vector $M_t = (M_{1,t}, M_{2,t}, \dots, M_{\bar{k},t})$ therefore takes $d = b_m^{\bar{k}}$ possible values $m^1, \dots, m^d \in \mathbb{R}_+^{\bar{k}}$. The dynamics of the Markov chain M_t are characterized by the transition matrix $A = (a_{i,j})_{1 \leq i,j \leq d}$ with components $a_{ij} = \mathbb{P}(M_{t+1} = m^j | M_t = m^i)$.⁵

Conditional on the volatility state M_t , the return x_t is Gaussian with density $f_{x_t}(x | M_t = m^i) = [\sigma g(m^i)]^{-1} n[x/\sigma g(m^i)]$, where $n(\cdot)$ denotes the density of a standard normal. The econometrician does not directly observe M_t , but can compute the conditional probabilities

$$\Pi_t^j \equiv \mathbb{P}(M_t = m^j | x_1, \dots, x_t)$$

over the unobserved states m^1, \dots, m^d . We can stack these probabilities in the row vector $\Pi_t = (\Pi_t^1, \dots, \Pi_t^d) \in \mathbb{R}_+^d$. Letting $\iota = (1, \dots, 1) \in \mathbb{R}^d$, we know that $\Pi_t \iota = 1$.

⁵ We note that $a_{ij} = \prod_{k=1}^{\bar{k}} [(1 - \gamma_k) 1_{\{m_k^i = m_k^j\}} + \gamma_k \mathbb{P}(M = m_k^j)]$, where m_k^i denotes the m th component of vector m^i , and $1_{\{m_k^i = m_k^j\}}$ is the dummy variable equal to one if $m_k^i = m_k^j$, and zero otherwise. In Calvet and Fisher (2001), the transition matrix differs because an innovation to a lower-frequency multiplier causes switching in all higher-frequency multipliers. Here we assume that arrival times are independent across frequencies.

The conditional probability vector Π_t is computed recursively. By Bayes' rule, Π_t can be expressed as a function of the previous belief Π_{t-1} and the innovation x_t :

$$\Pi_t = \frac{\omega(x_t) * (\Pi_{t-1}A)}{[\omega(x) * (\Pi_{t-1}A)]_t}, \quad (3)$$

where $x * y$ denotes the Hadamard product (x_1y_1, \dots, x_dy_d) for any $x, y \in \mathbb{R}^d$, and

$$\omega(x_t) = \left[\frac{n(x_t/\sigma g(m^1))}{\sigma g(m^1)}, \dots, \frac{n(x_t/\sigma g(m^d))}{\sigma g(m^d)} \right].$$

These results are familiar in regime switching models.⁶ In empirical applications, the initial vector Π_0 is chosen to be the ergodic distribution of the Markov process. Since the multipliers $(M_{1,1}, \dots, M_{\bar{k},1})$ are independent, the components of Π_0 satisfy $\prod_0^j = \prod_{t=1}^k \mathbb{P}(M = m_t^j)$ for all j .

2.2 Closed-Form Likelihood

Having solved the conditioning problem, we easily check that the log-likelihood function is

$$\ln L(x_1, \dots, x_T; \psi) = \sum_{t=1}^T \ln \left[\omega(x_t) \cdot \left(\Pi_{t-1}A \right) \right]. \quad (4)$$

For a fixed \bar{k} , we know that the ML estimator is consistent and asymptotically efficient as $T \rightarrow \infty$. Unlike standard stochastic volatility models [e.g. Ghysels, Harvey, and Renault (1996)], MSM thus has a closed-form likelihood. The parsimonious parameterization of the transition matrix represents an important difference between MSM and standard Markov switching models. This allows us to estimate MSM with reasonable precision even under a very large state space. While the expectation maximization (EM) algorithm [Hamilton (1990)] is not directly applicable with constrained transition probabilities, we now show that numerical optimization of the likelihood function produces good results.

2.3 Small-Sample Properties

The small-sample properties of the ML estimator are assessed by Monte Carlo simulations. We first examine parameter estimation when the specification has a fixed number \bar{k} of volatility components. We then analyze the choice of \bar{k} , which can be viewed as model selection. All simulations use binomial specifications where the multiplier M takes values m_0 and $2 - m_0$ with equal probability.

To evaluate the finite-sample properties of ML estimation, we use specifications with $\bar{k} = 8$ volatility components, which is representative of models performing well in later empirical sections. The simulation requires four parameters: the binomial value m_0 , the unconditional standard deviation σ , the frequency growth rate b , and the high-frequency switching probability $\gamma_{\bar{k}}$. The unconditional

⁶ See Hamilton (1994, chap. 22) for a derivation of Equations (3) and (4).

standard deviation is a simple scale factor, which is normalized to unity: $\sigma = 1$. Remaining parameters are set to values consistent with the empirical results. Specifically, all simulations use $b = 3$ and $\gamma_k = 0.95$. The binomial parameter takes one of three values: $m_0 \in \{1.3, 1.4, 1.5\}$. The sample length T belongs to $\{T_1, T_2, T_3\}$, where $T_1 = 2500$, $T_2 = 5000$, and $T_3 = 10,000$.

Table 1, panel A, reports the Monte Carlo results. Each of the nine columns corresponds to a different combination of the parameter m_0 and the sample length T . For each column, we simulate 400 independent sample paths of the corresponding model and sample length. ML estimation then provides a set of parameter estimates and asymptotic standard errors for each path.⁷ The table has four rows corresponding to each parameter. The first row gives the average point estimate over the simulated paths. The second row is the standard error of these point estimates, or the finite sample standard error (FSSE). The third row gives the root mean squared error (RMSE) of the parameter estimates relative to the true parameter values. Finally, the average asymptotic standard error (AASE) gives the average over the 400 simulations of the asymptotic standard errors calculated from the information matrix. As sample size becomes large, we expect that the AASE and the FSSE become close.

Maximum-likelihood estimation produces reasonable results. For m_0 , σ , and b , the biases are small and become negligible as sample size increases. The parameter m_0 has a low standard error relative to its size and is thus well identified, which is important because m_0 largely determines fluctuations in volatility. By contrast, the unconditional standard deviation σ has standard errors that, although declining as expected with sample length, are roughly 10% of the true parameter value. This result is consistent with the low-frequency variations contained in MSM, which create considerable uncertainty about long-run averages. We view this property as a strength of our model, which encourages the econometrician using MSM to be cautious about the long-run state of the economy. The parameter b shows a moderate degree of uncertainty about the spacing between frequencies. Finally, the high-frequency switching probability γ_k is the only parameter that shows more than a small bias. This disappears quickly as sample size increases, and the standard errors are not generally large. Overall the Monte Carlo simulations show that ML estimation produces reliable results given the sample sizes considered in subsequent sections. We also note that the convenience and efficiency of ML offers significant advantages relative to previous moment-based estimators for multifractal processes.⁸

We next investigate by Monte Carlo whether the integer \bar{k} can be reasonably selected by ML. The simulations assume $\bar{k} = 5$ volatility components and parameter values $(m_0, \sigma, b, \gamma_k) = (1.5, 0.5, 8, 0.75)$ representative of later empirical

⁷ We start the optimizations at the true parameter values and iterate to convergence once. Preliminary work considered searching for multiple local optima, and although these occasionally exist, they do not significantly affect the means or standard deviations of the reported Monte Carlo results.

⁸ Early drafts of the article experimented with other computational methods, including Simulated Method of Moments (SMM) and Efficient Method of Moments (EMM), which were found to be substantially less efficient in finite sample than the ML estimator presented here.

Table 1 Monte Carlo MLE results.

Panel A. Parameter estimation									
	$m_0 = 1.3$			$m_0 = 1.4$			$m_0 = 1.5$		
	T_1	T_2	T_3	T_1	T_2	T_3	T_1	T_2	T_3
\bar{m}_{sim}	1.288	1.293	1.297	1.392	1.393	1.397	1.494	1.494	1.497
FSSE	(0.026)	(0.018)	(0.012)	(0.031)	(0.019)	(0.015)	(0.032)	(0.025)	(0.015)
RMSE	(0.029)	(0.019)	(0.012)	(0.032)	(0.021)	(0.015)	(0.033)	(0.025)	(0.016)
AASE	(0.019)	(0.013)	(0.010)	(0.018)	(0.014)	(0.011)	(0.019)	(0.014)	(0.011)
$\bar{\sigma}_{sim}$	1.014	1.004	0.999	1.031	1.011	1.005	1.026	1.017	1.006
FSSE	(0.167)	(0.102)	(0.073)	(0.221)	(0.147)	(0.091)	(0.255)	(0.224)	(0.105)
RMSE	(0.167)	(0.102)	(0.073)	(0.223)	(0.148)	(0.092)	(0.256)	(0.225)	(0.105)
AASE	(0.072)	(0.056)	(0.046)	(0.091)	(0.075)	(0.057)	(0.111)	(0.088)	(0.067)
$\bar{\gamma}_{sim}$	0.867	0.908	0.934	0.907	0.935	0.940	0.924	0.938	0.944
FSSE	(0.164)	(0.113)	(0.068)	(0.111)	(0.069)	(0.047)	(0.087)	(0.055)	(0.037)
RMSE	(0.184)	(0.120)	(0.070)	(0.119)	(0.070)	(0.048)	(0.091)	(0.056)	(0.037)
AASE	(0.181)	(0.114)	(0.071)	(0.110)	(0.068)	(0.048)	(0.078)	(0.052)	(0.036)
\bar{b}_{sim}	2.853	2.942	2.988	3.052	2.938	2.973	3.054	2.987	2.979
FSSE	(0.934)	(0.670)	(0.416)	(0.963)	(0.480)	(0.363)	(0.735)	(0.565)	(0.316)
RMSE	(0.945)	(0.673)	(0.416)	(0.964)	(0.484)	(0.364)	(0.737)	(0.565)	(0.317)
AASE	(0.677)	(0.471)	(0.331)	(0.599)	(0.367)	(0.264)	(0.476)	(0.323)	(0.225)
Panel B. Volatility components ($\bar{k} = 5$)									
	$\bar{k}' = 1$	2	3	4	5	6	7	8	
Times chosen (max = 400)	0	0	2	81	299	13	5	0	
Mean[$\ln L(\bar{k}') - \ln L(\bar{k} = 5)$]	-425.30	-118.44	-42.53	-13.07	0	-0.89	-1.17	-0.97	
	(5.37)	(1.70)	(1.15)	(0.54)	—	(0.06)	(0.09)	(0.02)	

Panel A is based on $J = 400$ simulated paths for each column. All simulations are based on a multifractal process with $\bar{k} = 8$. The columns are distinguished by combinations of $m_0 \in \{1.3, 1.4, 1.5\}$ and sample lengths of $T_1 = 2500$, $T_2 = 5000$, and $T_3 = 10,000$. Parameters that are fixed across all simulations are $\sigma = 1$, $\gamma_k = 0.95$, and $b = 3$. These parameters provide a reasonable approximation to values that are estimated on exchange rate data in future sections of the article. In each set of four rows, the first row is the average MLE parameter value over the J simulated paths. In the remaining rows, FSSE denotes finite sample standard error, RMSE the root mean squared error, and AASE the average asymptotic standard error. To derive AASE, the asymptotic variance is calculated for each $j \in \{1, \dots, J\}$ path from the inverse of the information matrix and the average is taken over the J simulations. Panel B is based on 400 simulated paths of length 10,000 from a process with five volatility components and parameters $(m_0, \sigma, b, \gamma_k) = (1.5, 0.5, 8, 0.75)$. For each path, parameters are estimated for $\text{MSM}(\bar{k}')$, $\bar{k}' \in \{1, \dots, 8\}$. The first row reports the number of times each value of \bar{k}' has highest likelihood. The second row shows the average likelihood difference relative to the true process with $\bar{k} = 5$. The final row gives in parentheses simulation standard errors, which are all small relative to the reported mean differences.

results. We simulate 400 paths of the process with $T=10,000$ observations. For each simulated path we compute the MLs corresponding to the specifications $\text{MSM}(\bar{k}')$, $\bar{k}' \in \{1, \dots, 8\}$, and select the integer \bar{k}' yielding the highest likelihood. Table 1, panel B, reports the number of times each \bar{k}' is chosen. We observe that the

true value $\bar{k}=5$ is selected in most cases. Of the remaining possibilities, only $\bar{k}'=4$ is selected more than a nominal number of times. We explain this by observing that the most persistent volatility component has a frequency comparable to the sample length. In some sample paths, few or no shifts occur in the most persistent component, and the process $\text{MSM}(\bar{k}'=4)$ then produces a higher likelihood. This result reflects the general fact that low-frequency events may be difficult to detect in a finite sample (the peso problem). As a result, the estimation method on average slightly underpredicts the number of components, but errors in \bar{k} of more than one component are very unlikely.

We also report in Table 1, panel B, the mean difference in optimized likelihoods between the $\text{MSM}(\bar{k}')$ and true $\text{MSM}(\bar{k}=5)$ specifications.⁹ All of the mean differences are negative, again confirming that the estimation method correctly identifies the true value of \bar{k} . Interestingly, having too few frequencies causes on average a larger penalty in the likelihood function than having too many. In empirical work we should thus expect a rapid increase in the likelihood as components are added, followed by flattening and slow decline. Nonetheless, a specification with too many components is very unlikely to give a higher likelihood. We can thus be confident that the selection of a large \bar{k} in the empirical work will not arise from estimation problems but should instead be attributed to heterogeneity in volatility shocks. Overall we conclude that ML estimation of the parameters and the frequency number \bar{k} produces reliable results in finite samples.

3 EMPIRICAL RESULTS

Using a binomial specification for the multiplier M , we apply ML estimation to four exchange rate series and obtain preferred specifications with a large number of volatility frequencies.

3.1 Exchange Rate Data

The empirical analysis uses daily exchange rate data for the deutsche mark (DEM), Japanese yen (JPY), British pound (GBP), and Canadian dollar (CAD), all against the U.S. dollar. The data consist of daily prices reported at noon by the Federal Reserve Bank of New York.¹⁰ The fixed exchange rate system broke down in early 1973, and the DEM, JPY, and GBP series accordingly begin on June 1, 1973. The CAD series starts a year later (June 1, 1974) because the Canadian currency was held essentially at parity with the U.S. dollar for several months after the demise of Bretton Woods. The deutsche mark was replaced by the euro at the beginning of 1999. The DEM data thus ends on December 31, 1998, while the other three series run until June 30, 2002. Overall the series contains 6420 observations for the deutsche mark, 7049 observations for the Canadian dollar, and 7299 observations for the yen and the pound.

⁹ Standard errors due to simulation are reported in parentheses and show that the simulated means are relatively accurate.

¹⁰ More specifically, the data consist of buying rates for wire transfers at 12:00 P.M. Eastern time.




Figure 2 Exchange rate data. This figure shows the daily log-price differences of the four exchange rate series.

Figure 2 plots the daily returns of each series. Consistent with earlier studies, we observe apparent volatility clustering on a range of frequencies. For each series, we compute in Table 2 the standard deviation of returns over the entire sample and over four subsamples of equal length. We observe that the sample

Table 2 FX return variability.

		Standard deviation of returns				
		Entire sample	By subperiod			
			1	2	3	4
DEM	0.664		0.587	0.716	0.708	0.635
JPY	0.657		0.545	0.640	0.646	0.775
GBP	0.607		0.486	0.724	0.699	0.473
CAD	0.274		0.220	0.255	0.284	0.327

For each dataset, the first column shows the standard deviation of returns over the entire sample period. In the next four columns, each data series is broken into quarters and the same statistic is calculated for each subperiod. The results show that the variability of return variance is substantial even at very low frequencies.

standard deviation varies quite substantially across subperiods, which is consistent with the low-frequency regime shifts in MSM.

3.2 ML Estimation Results

Table 3 reports ML estimation results for all four currencies. The columns of the table correspond to the number of frequencies \bar{k} varying from 1 to 10. The first column is a standard Markov switching model with only two possible values for volatility. As \bar{k} increases, the number of states increases at the rate $2^{\bar{k}}$. There are thus more than 1000 states when $\bar{k} = 10$.

We begin by examining the DEM data. The multiplier parameter \hat{m}_0 tends to decline with \bar{k} : with a larger number of components, less variability is required in each M_{kt} to match the fluctuations in volatility exhibited by the data. The estimates of $\hat{\sigma}$ vary across \bar{k} with no particular pattern. Standard errors of $\hat{\sigma}$ increase with \bar{k} , consistent with the idea that long-run averages are difficult to identify in models permitting long volatility cycles. We next examine the frequency parameters $\hat{\gamma}_{\bar{k}}$ and \hat{b} . When $\bar{k} = 1$, the single multiplier has a duration slightly less than two weeks. As \bar{k} increases, the switching probability of the highest-frequency multiplier increases until a switch occurs about once a day for large \bar{k} . At the same time, the estimate \hat{b} decreases steadily with \bar{k} . When $\bar{k} = 10$, we infer from Equation (2) that the lowest-frequency multiplier has a duration approximately equal to 10 years, or about one-third the sample size. Thus, as \bar{k} increases, the range of frequencies spreads out while the spacing between frequencies becomes tighter.

The other currencies generate parameter estimates with similar properties. In all cases, \hat{m}_0 tends to decrease with \bar{k} . The values of \hat{m}_0 , and thus the importance of stochastic volatility, are largest for JPY and GBP and smallest for CAD. Variability across \bar{k} in the estimates of $\hat{\sigma}$ is also greatest for JPY and GBP and least for CAD. As \bar{k} increases, the most transitory multiplier switches more often and the spacing between frequencies becomes tighter for all currencies. The spacing is widest for JPY and GBP and tightest for CAD. We observe correspondingly that the most

persistent multiplier has the longest duration for JPY at approximately three times the sample size and the smallest for CAD at approximately one-tenth the sample size. All standard errors are roughly consistent with the magnitudes obtained in our Monte Carlo simulations.

For large values of \bar{k} , the estimated MSM(\bar{k}) processes generate substantial outliers despite having finite moments of every order, as is now shown. For each currency we use the estimated process with $\bar{k}=10$ frequencies to generate 10,000 paths of the same length as the data and compute a Hill tail index α for each simulated path. Basing the index on 100 order statistics, the empirical tail index and the average α in the simulated samples are equal to 4.74 and 4.34 (DEM), 3.91 and 3.75 (JPY), 4.59 and 4.03 (GBP), and 4.40 and 4.79 (CAD), respectively. Furthermore, for all currencies, we cannot at the 10% level reject equality of the simulated and empirical tail statistics. This result is caused by the high-frequency variations in volatility in the estimated models. The distribution of returns in MSM is a mixture of Gaussians, which has finite moments of every order. With the highest-frequency multipliers taking new values almost daily, this mixture appears to be more than sufficient to capture the tail characteristics of the data, even in a sample containing 30 years of daily observations.

Finally, we examine the behavior of the log-likelihood function as the number of frequencies \bar{k} increases from 1 to 10. For each currency, the likelihood goes up substantially at low \bar{k} , and in most cases continues to increase at a decreasing rate. The only exception to the monotonic increase in likelihood occurs in the DEM series, for which the likelihood reaches a peak at $\bar{k}=7$. In all other cases the likelihood reaches a maximum at $\bar{k}=10$. This behavior of the likelihood confirms one of the main premises of the multifractal approach: fluctuations in volatility occur with heterogeneous degrees of persistence, and explicitly incorporating a larger number of frequencies results in a better fit.

3.3 Model Selection

We now examine the statistical significance of the differences in likelihoods across estimated MSM(\bar{k}) processes. Consider two models MSM(\bar{k}) and MSM(\bar{k}'), $\bar{k} \neq \bar{k}'$, with respective densities f and g . The processes are nonnested and have log-likelihood difference

$$\sqrt{T}(\ln L_T^f - \ln L_T^g) = \frac{1}{\sqrt{T}} \sum_{t=1}^T \ln \frac{f(x_t | x_1, \dots, x_{t-1})}{g(x_t | x_1, \dots, x_{t-1})}.$$

Consider the null hypothesis that the models have identical unconditional expected log-likelihoods. When the observations x_t are i.i.d., Vuong (1989) shows that the difference $\ln L_T^f - \ln L_T^g$ is asymptotically normal under the null.¹¹ In addition, the variance of this difference is consistently estimated by the sample variance of the addends $\ln[f(x_t | x_1, \dots, x_{t-1})/g(x_t | x_1, \dots, x_{t-1})]$. Since the observations $\{x_t\}$ are

¹¹ See the appendix for a more detailed discussion of Vuong tests.

Table 3 Maximum likelihood results.

	$\bar{k} = 1$	2	3	4	5	6	7	8	9	10
Deutsche mark/U.S. dollar										
\hat{m}_0	1.654 (0.013)	1.590 (0.012)	1.555 (0.013)	1.492 (0.013)	1.462 (0.012)	1.413 (0.013)	1.380 (0.012)	1.353 (0.011)	1.351 (0.013)	1.326 (0.015)
$\hat{\sigma}$	0.682 (0.012)	0.651 (0.018)	0.600 (0.014)	0.572 (0.016)	0.512 (0.018)	0.538 (0.026)	0.547 (0.021)	0.550 (0.025)	0.674 (0.035)	0.643 (0.073)
$\hat{\gamma}_{\bar{k}}$	0.075 (0.011)	0.107 (0.022)	0.672 (0.151)	0.714 (0.096)	0.751 (0.106)	0.858 (0.128)	0.932 (0.071)	0.974 (0.042)	0.966 (0.065)	0.959 (0.066)
\hat{b}	–	8.01 (2.58)	21.91 (7.30)	10.42 (1.92)	7.89 (1.31)	5.16 (0.76)	4.12 (0.48)	3.38 (0.36)	3.29 (0.47)	2.70 (0.36)
$\ln L$	–5920.86	–5782.96	–5731.78	–5715.31	–5708.25	–5706.91	5704.48	–5704.77	–5704.86	–5705.09
Japanese yen/U.S. dollar										
\hat{m}_0	1.797 (0.011)	1.782 (0.009)	1.693 (0.010)	1.654 (0.010)	1.640 (0.010)	1.573 (0.010)	1.565 (0.010)	1.513 (0.010)	1.475 (0.010)	1.448 (0.011)
$\hat{\sigma}$	0.630 (0.011)	0.538 (0.009)	0.566 (0.017)	0.462 (0.013)	0.709 (0.023)	0.642 (0.023)	0.518 (0.018)	0.514 (0.020)	0.486 (0.026)	0.461 (0.036)
$\hat{\gamma}_{\bar{k}}$	0.199 (0.019)	0.345 (0.033)	0.312 (0.054)	0.697 (0.080)	0.778 (0.076)	0.899 (0.060)	0.897 (0.057)	0.975 (0.034)	0.995 (0.010)	0.998 (0.006)
\hat{b}	–	134.20 (48.27)	12.46 (2.18)	15.58 (2.67)	16.03 (2.67)	8.07 (1.03)	7.46 (0.89)	5.65 (0.78)	4.43 (0.53)	3.76 (0.45)
$\ln L$	–6451.80	–6102.18	–5959.72	–5900.67	–5882.93	–5871.35	–5867.88	–5863.20	–5863.01	–5862.68
British pound/U.S. Dollar										
\hat{m}_0	1.716 (0.012)	1.671 (0.011)	1.648 (0.011)	1.609 (0.011)	1.579 (0.011)	1.534 (0.012)	1.503 (0.012)	1.461 (0.011)	1.428 (0.011)	1.403 (0.009)
$\hat{\sigma}$	0.609 (0.009)	0.590 (0.011)	0.513 (0.016)	0.467 (0.016)	0.421 (0.017)	0.468 (0.019)	0.389 (0.014)	0.384 (0.015)	0.374 (0.022)	0.370 (0.022)
$\hat{\gamma}_{\bar{k}}$	0.110	0.222	0.278	0.645	0.637	0.784	0.811	0.958	0.964	0.982

\hat{b}	—	(0.017)	(0.034)	(0.052)	(0.080)	(0.075)	(0.078)	(0.083)	(0.052)	(0.043)	(0.031)
		19.90	14.29	12.51	11.02	8.32	6.72	5.23	4.08	3.45	
		(5.19)	(2.58)	(2.00)	(1.74)	(1.15)	(0.91)	(0.69)	(0.41)	(0.32)	
$\ln L$	−5960.18	−5724.37	−5622.73	−5570.02	−5537.80	−5523.64	−5516.89	−5515.37	−5515.28	−5514.94	
Canadian dollar/U.S. dollar											
\hat{m}_0	1.646	1.556	1.474	1.435	1.386	1.374	1.338	1.319	1.296	1.278	
	(0.012)	(0.012)	(0.014)	(0.015)	(0.012)	(0.013)	(0.012)	(0.016)	(0.013)	(0.012)	
$\hat{\sigma}$	0.280	0.278	0.293	0.263	0.251	0.295	0.282	0.262	0.259	0.262	
	(0.005)	(0.006)	(0.014)	(0.009)	(0.010)	(0.011)	(0.013)	(0.017)	(0.015)	(0.021)	
$\hat{\gamma}_k$	0.064	0.109	0.129	0.171	0.441	0.524	0.593	0.594	0.631	0.644	
	(0.009)	(0.016)	(0.040)	(0.062)	(0.153)	(0.128)	(0.145)	(0.151)	(0.155)	(0.158)	
\hat{b}	—	10.92	4.76	3.95	4.02	4.08	3.11	2.72	2.35	2.11	
		(3.12)	(1.15)	(0.83)	(0.76)	(0.58)	(0.39)	(0.39)	(0.25)	(0.18)	
$\ln L$	−271.01	−129.80	−105.16	−91.32	−88.41	−84.73	−84.03	−83.40	−83.06	−83.00	

This table shows ML estimation results for the binomial multifractal model for all four exchange rate series. Columns correspond to the number of frequencies \bar{k} in the estimated model. The likelihood function increases monotonically in the number of volatility frequencies for all datasets except DEM, which obtains a maximum at $\bar{k} = 7$. Asymptotic standard errors are in parentheses.

Table 4 Multifractal model selection.

	$\bar{k} = 1$	2	3	4	5	6	7	8	9
Panel A. Vuong (1989) test									
Mark	-8.655 (0.000)	-5.523 (0.000)	-2.972 (0.001)	-1.858 (0.032)	-0.688 (0.246)	-0.733 (0.232)	0.341 (0.633)	0.204 (0.581)	0.337 (0.632)
Yen	-13.067 (0.000)	-8.406 (0.000)	-5.342 (0.000)	-3.154 (0.001)	-2.156 (0.016)	-1.192 (0.117)	-1.108 (0.134)	-0.180 (0.429)	-0.162 (0.436)
Pound	-11.810 (0.000)	-8.337 (0.000)	-6.267 (0.000)	-4.360 (0.000)	-2.984 (0.001)	-1.334 (0.089)	-0.408 (0.342)	-0.149 (0.441)	-0.236 (0.407)
Canada	-8.475 (0.000)	-4.421 (0.000)	-3.289 (0.000)	-1.795 (0.036)	-2.108 (0.017)	-0.862 (0.194)	-0.825 (0.205)	-0.472 (0.318)	-0.158 (0.437)
Panel B. HAC-adjusted Vuong test									
Mark	-4.285 (0.000)	-3.033 (0.001)	-1.683 (0.046)	-1.101 (0.135)	-0.402 (0.344)	-0.424 (0.336)	0.197 (0.578)	0.120 (0.548)	0.194 (0.577)
Yen	-5.219 (0.000)	-4.262 (0.000)	-2.865 (0.002)	-1.645 (0.050)	-1.224 (0.111)	-0.648 (0.259)	-0.663 (0.254)	-0.105 (0.458)	-0.098 (0.461)
Pound	-3.788 (0.000)	-2.804 (0.003)	-2.803 (0.003)	-2.195 (0.014)	-1.759 (0.039)	-0.779 (0.218)	-0.242 (0.404)	-0.088 (0.465)	-0.137 (0.446)
Canada	-4.237 (0.000)	-2.383 (0.009)	-1.789 (0.037)	-1.019 (0.154)	-1.150 (0.125)	-0.480 (0.316)	-0.445 (0.328)	-0.276 (0.391)	-0.091 (0.464)

This table reports t -ratios and one-sided p -values for the log-likelihood difference of the model in each column against the multifractal with 10 frequencies. Panel A uses the Vuong (1989) methodology and panel B adjusts for heteroskedasticity and autocorrelation using Newey and West (1987, 1994). A low p -value indicates that the corresponding model would be rejected in favor of the multifractal with 10 frequencies.

typically not i.i.d. in financial applications, we construct in the appendix a HAC-adjusted version of the Vuong test. Our discussion is a simplified version of the broader approach recently proposed by Rivers and Vuong (2002).

For each $\bar{k} \in \{1, \dots, 9\}$, we test in Table 4 the null hypothesis that $\text{MSM}(\bar{k})$ and $\text{MSM}(10)$ fit the data equally well. Since HAC-adjusted tests tend to perform poorly in small samples [see, e.g., Andrews (1991), Andrews and Monahan (1992), and den Haan and Levin (1997)], we compute t -ratios and one-sided p -values using both the original and the HAC-adjusted methods. For $\bar{k} \in \{1, 2, 3\}$, the log-likelihood difference is significant at the 1% level in the nonadjusted case (Table 4, panel A) and at the 5% level in the HAC case (Table 4, panel B). This is strong evidence that $\text{MSM}(10)$ significantly outperforms models with one to three frequencies. For $\bar{k} \in \{4, 5\}$, we reject the null at the 5% (nonadjusted) and 20% (HAC-adjusted) levels in almost all cases. These results represent rather substantial evidence that $\text{MSM}(10)$ outperforms models with four or five frequencies. Lower significance levels are obtained for larger values of \bar{k} , and the overall conclusion is that the MSM model works better for larger \bar{k} . For this reason, and consistency in the remaining analysis, we henceforth focus on the $\text{MSM}(\bar{k}=10)$ process for all currencies.

4 COMPARISON WITH ALTERNATIVE MODELS

This section compares the multifractal model with GARCH(1,1), Markov switching GARCH (MS-GARCH), and fractionally integrated GARCH (FIGARCH). We show in subsections 4.1 and 4.2 that MSM outperforms GARCH and MS-GARCH both in- and out-of-sample. The multifractal is then compared with FIGARCH in subsection 4.3 to assess the connection between long memory and forecasting performance.

4.1 In-Sample Comparison

We consider alternative processes of the form $x_t = h_t^{1/2} e_t$, where h_t is the conditional variance of x_t at date $t - 1$, and $\{e_t\}$ are i.i.d. Student innovations with unit variance and ν degrees of freedom (d.f.). GARCH(1,1) assumes the recursion $h_{t+1} = \omega + \alpha \varepsilon_t^2 + \beta h_t$. MS-GARCH combines short-run autoregressive dynamics with low-frequency regime shifts [Gray (1996), Klaassen (2002)]. A latent state $s_t \in \{1, 2\}$ follows a first-order Markov process with transition probabilities $p_{ij} = \mathbb{P}(s_{t+1} = j | s_t = i)$. In every period the econometrician observes the return x_t but not the latent s_t . For $i = \{1, 2\}$, let $h_{t+1}(i) = \text{var}_t(x_{t+1} | s_{t+1} = i)$ denote the variance of x_{t+1} conditional on past returns $\{x_s\}_{s=1}^t$ and $s_{t+1} = i$. The quantity h_t is latent in every period, and the econometrician can similarly define $\mathbb{E}_t[h_t(s_t) | s_{t+1} = i]$, the expectation of h_t conditional on $s_{t+1} = i$ and past returns. Klaassen (2002) assumes the conditional dynamics:

$$h_{t+1}(i) = \omega_i + \alpha_i \varepsilon_t^2 + \beta_i \mathbb{E}_t[h_t(s_t) | s_{t+1} = i]. \quad (5)$$

The equation conditions volatility on a larger information set than the Gray specification: $h_{t+1}(i) = \omega_i + \alpha_i \varepsilon_t^2 + \beta_i \mathbb{E}_{t-1} h_t(s_t)$. We prefer Equation (5) for two reasons. First, Klaassen shows that his model provides better forecasts for three of the exchange rates considered in this article (DEM, JPY, and GBP), and attributes these improvements to finer conditioning. Second, the Klaassen version improves on Gray (1996) by permitting analytical multistep forecasting.

We report in Table 5 the ML estimates of GARCH and MS-GARCH. The coefficient $1/\nu$ is the inverse of the degrees of freedom in the Student distribution. This convenient renormalization has been frequently used in the literature [e.g., Bollerslev (1987)]. Each coefficient σ_i , $i = 1, 2$, represents the standard deviation of returns conditional on the volatility state: $\sigma_i^2 = \omega_i / (1 - \alpha_i - \beta_i)$. These coefficients are easier to interpret across models than the intercepts ω_i . As shown in Table 6, the multifractal has a higher likelihood than GARCH(1,1) for all exchange rates, even though both processes have the same number of parameters. Note that GARCH(1,1) approximately matches the likelihoods obtained by MSM with only three or four frequencies. The multifractal model thus gives an improved fit over GARCH(1,1) in-sample.

The MS-GARCH model uses nine parameters as compared to four with either GARCH or MSM. Thus, although MS-GARCH has higher likelihoods than either GARCH or MSM, we obtain a different ordering when using the Schwartz

Table 5 Alternative processes.

		Regime 1				Regime 2				
	$1/\nu$	σ_1	α_1	β_1	p_{11}	σ_2	α_2	β_2	p_{22}	$\ln L$
Deutsche mark/U.S. dollar										
GARCH	0.1929 (0.011)	1.5539 (0.405)	0.0879 (0.009)	0.9108 (0.009)						-5730.52
MS-GARCH	0.2041 (0.011)	1.0749 (0.288)	0.2048 (0.023)	0.7896 (0.024)	0.9998 (0.0003)	1.3145 (0.282)	0.0718 (0.010)	0.9241 (0.011)	0.9999 (0.0002)	-5694.78
Japanese yen/U.S. dollar										
GARCH	0.2290 (0.0002)	0.1638 (0.059)	0.0652 (0.006)	0.9348 (0.006)						-5965.07
MS-GARCH	0.2632 (0.012)	0.4443 (0.137)	0.3420 (0.040)	0.6500 (0.040)	0.9999 (0.0002)	0.9639 (0.121)	0.0650 (0.010)	0.9227 (0.013)	0.9999 (0.0002)	-5833.59
British pound/U.S. dollar										
GARCH	0.2007 (0.008)	0.2365 (0.070)	0.0681 (0.005)	0.9319 (0.005)						-5562.00
MS-GARCH	0.2202 (0.009)	0.8423 (0.013)	0.3653 (0.053)	0.6051 (0.056)	0.9860 (0.005)	0.9343 (0.012)	0.0587 (0.008)	0.9365 (0.008)	0.9986 (0.0003)	-5492.44
Canadian dollar/U.S. dollar										
GARCH	0.1528 (0.037)	0.3108 (0.008)	0.0810 (0.008)	0.9108 (0.010)						-96.03
MS-GARCH	0.1385 (0.011)	0.2046 (0.035)	0.0584 (0.009)	0.9361 (0.010)	0.9896 (0.004)	0.2972 (0.025)	0.2587 (0.074)	0.2925 (0.215)	0.9415 (0.023)	-73.51

This table shows ML estimation results for alternative processes for the four exchange rate series. Asymptotic standard errors are in parentheses. For the GARCH(1,1) model, the parameter estimates for JPY/USD and GBP/USD are on the boundary of the restriction $\alpha + \beta \leq 1 - \epsilon$, where $\epsilon = 10^{-5}$.

Table 6 In-sample model comparison.

				BIC p -value vs. multifractal	
	No. of parameters	$\ln L$	BIC	Vuong (1989)	HAC Adj
Deutsche mark/U.S. dollar					
Binomial multifractal	4	−5705.09	1.7830		
GARCH	4	−5730.52	1.7910	0.005	0.071
MS-GARCH	9	−5694.78	1.7866	0.140	0.248
Japanese yen/U.S. dollar					
Binomial multifractal	4	−5862.68	1.6115		
GARCH	4	−5965.07	1.6396	0.000	0.008
MS-GARCH	9	−5833.59	1.6097	0.619	0.572
British pound/U.S. dollar					
Binomial multifractal	4	−5514.94	1.5162		
GARCH	4	−5562.00	1.5291	0.004	0.070
MS-GARCH	9	−5492.44	1.5162	0.505	0.503
Canadian dollar/U.S. dollar					
Binomial multifractal	4	−83.00	0.0286		
GARCH	4	−96.03	0.0323	0.072	0.200
MS-GARCH	9	−73.51	0.0322	0.092	0.235

This table summarizes information about in-sample goodness-of-fit for the three models. The Bayesian information criterion is given by $BIC = T^{-1}(-2 \ln L + NP \ln T)$. The sample lengths are 6419 for DEM/USD, 7298 for JPY/USD and GBP/USD, and 7048 for CAD/USD. The last two columns give p -values from a test that the corresponding model dominates the multifractal model by the BIC. The first value uses the Vuong (1989) methodology and the second value adjusts the test for heteroscedasticity and autocorrelation. A low p -value indicates that the corresponding model would be rejected in favor of the multifractal model.

BIC criterion. In particular, the multifractal model is then indistinguishable from MS-GARCH in the GBP data, and is preferred for DEM and CAD.

As suggested by Vuong (1989), we evaluate the statistical significance of BIC differences. The last two columns of Table 6 test the GARCH and MS-GARCH models against MSM under this metric.¹² We again give p -values for both the original version of the test as well as a HAC-adjusted variant. Under the original version, the in-sample performance of MSM over GARCH(1,1) is highly significant for DEM, JPY, and GBP, and somewhat significant for CAD. The HAC adjustments produce analogous, but slightly weaker results. For DEM and CAD, there is some evidence that the multifractal model is a better performer than MS-GARCH, but the significance is marginal at best. Overall the in-sample analysis suggests that the multifractal matches the performance MS-GARCH and significantly outperforms GARCH(1,1).

¹² Note that a BIC test of GARCH against the multifractal model is identical to a likelihood test, since both have the same number of parameters.

4.2 Out-of-Sample Forecasts

We now investigate the out-of-sample performance of the competing models over forecasting horizons ranging from 1 to 50 days. For each currency we estimate the three processes on the beginning of the series and use the last 12 years of data (or approximately half the sample) for out-of-sample comparison.

Table 7 reports results of one-day forecasts. The first two columns correspond to the coefficients γ_0 and γ_1 from the Mincer-Zarnowitz ordinary least squares (OLS) regressions $x_{t+1}^2 = \gamma_0 + \gamma_1 \mathbb{E}_t x_{t+1}^2 + u_t$ of squared returns on a constant and one-day forecasts. These regressions are common in the financial econometrics literature [e.g., Pagan and Schwert (1990); West and Cho (1995); Andersen, Bollerslev, and Meddahi (2002)], and unbiased forecasts would imply $\gamma_0 = 0$ and $\gamma_1 = 1$. We adjust the standard errors of γ_0 and γ_1 for parameter uncertainty as in West and McCracken (1998), and for HAC effects using the weighting and lag selection methodology of Newey and West (1987, 1994).

With the multifractal, the estimated intercept $\hat{\gamma}_0$ is slightly positive and the slope $\hat{\gamma}_1$ is slightly less than one for all currencies. These small biases, however, are not statistically significant. In particular, the hypothesis $\gamma_0 = 0$ is accepted at the 5% confidence level for all currencies, and $\gamma_1 = 1$ is accepted at the 5% level for JPY and CAD and at the 1% level for DEM and GBP. The Mincer-Zarnowitz regressions thus show little evidence of bias in MSM forecasts.

The point estimates $\hat{\gamma}_0$ and $\hat{\gamma}_1$ are slightly poorer with GARCH(1,1) than with the multifractal. All intercepts are more positive, and the slopes are further away from one for three currencies. The biases are also statistically significant. The hypotheses $\gamma_0 = 0$ and $\gamma_1 = 1$ are rejected at the 5% level in seven of eight cases. Since $0 < \hat{\gamma}_1 < 1$, these results suggest that GARCH forecasts are too variable and can be improved by the linear smoothing $\hat{\gamma}_0 + \hat{\gamma}_1 \mathbb{E}_t x_{t+1}^2$. In contrast, Markov switching GARCH improves on the out-of-sample performance of GARCH(1,1). We accept that $\gamma_0 = 0$ at the 5% confidence level for all currencies, and that $\gamma_1 = 1$ at the 1% level for DEM, GBP, and CAD. Furthermore, the regression estimates are best with MS-GARCH for two currencies (DEM and GBP) and with the multifractal for the other two. MSM and MS-GARCH thus seem to perform quite similarly out-of-sample at the one-day horizon. We also report in Table 7 two standard measures of goodness-of-fit: the mean squared error (MSE) and the restricted R^2 coefficient.¹³ The multifractal produces the best forecasting R^2 for DEM and JPY. On the other hand, GARCH produces better results for the GBP and MS-GARCH for the CAD. To summarize the one-day forecast results, the multifractal model appears to slightly dominate GARCH(1,1) and to give results comparable to MS-GARCH.

Multistep forecasts provide stronger empirical differences between the three models. We report in Table 8 the results for 20-day forecasts, which are representative of longer horizons. Since the data contain only business days, this frequency corresponds to about a month of calendar time. Following Andersen and Bollerslev

¹³ The MSE quantifies forecast errors in the out-of-sample period: $L^{-1} \sum_{t=T-L+1}^T (x_t^2 - \mathbb{E}_{t-1} x_t^2)^2$. The coefficient of determination is defined by $R^2 = 1 - \text{MSE}/\text{TSS}$, where TSS is the out-of-sample variance of squared returns: $\text{TSS} = L^{-1} \sum_{t=T-L+1}^T (x_t^2 - \sum_{i=T-L+1}^T x_i^2 / L)^2$.

Table 7 One-day forecasts.

	Mincer-Zarnowitz		Restricted $\gamma_0 = 0, \gamma_1 = 1$	
	γ_0	γ_1	MSE	R^2
Deutsche mark/U.S. dollar				
Binomial multifractal	0.098 (0.072)	0.703 (0.126)	0.7263	0.041
GARCH	0.153 (0.061)	0.622 (0.105)	0.7304	0.035
MS-GARCH	0.042 (0.080)	0.740 (0.130)	0.7296	0.037
Japanese yen/U.S. dollar				
Binomial multifractal	0.028 (0.090)	0.772 (0.117)	1.6053	0.053
GARCH	0.172 (0.075)	0.668 (0.105)	1.6137	0.048
MS-GARCH	0.080 (0.084)	0.709 (0.109)	1.6141	0.048
British pound/U.S. dollar				
Binomial multifractal	0.053 (0.049)	0.715 (0.100)	0.5081	0.057
GARCH	0.085 (0.044)	0.751 (0.098)	0.4980	0.076
MS-GARCH	0.017 (0.051)	0.814 (0.108)	0.4997	0.072
Canadian dollar/U.S. dollar				
Binomial multifractal	0.015 (0.016)	0.905 (0.156)	0.0345	0.051
GARCH	0.033 (0.012)	0.679 (0.111)	0.0348	0.042
MS-GARCH	0.025 (0.013)	0.785 (0.124)	0.0344	0.055

This table gives out-of-sample forecasting results for the three models. The first two columns correspond to parameter estimates from the Mincer-Zarnowitz OLS regression $e_{t+1}^2 = \gamma_0 + \gamma_1 E_t(e_{t+1}^2) + u_t$. For an unbiased forecast we expect $\gamma_0 = 0$ and $\gamma_1 = 1$. Asymptotic standard errors in parentheses are corrected for heteroscedasticity and autocorrelation using the method of Newey and West (1987, 1994) and for parameter uncertainty using the method of West and McCracken (1998). MSE is the mean squared forecast error, and R^2 is one less the MSE divided by the sum of squared demeaned squared returns in the out-of-sample period.

(1998) and Klaassen (2002), the dependent variable is the sum of squared daily returns $\sum_{s=t}^{t+19} x_s^2$ over the 20-day period. Because the average size of returns increases with the sampling interval, the estimated intercepts $\hat{\gamma}_0$ are larger in Table 8 than in Table 7. For each currency, the multifractal produces point estimates of γ_0 and γ_1 that are closest to their preferred values. We also accept the hypotheses

Table 8 Twenty-day forecasts.

	Mincer-Zarnowitz		Restricted	
	γ_0	γ_1	$\gamma_0 = 0,$ MSE	$\gamma_1 = 1$ R^2
Deutsche mark/U.S. dollar				
Binomial multifractal	1.749 (1.649)	0.706 (0.150)	37.12	0.135
GARCH	4.474 (1.108)	0.443 (0.092)	49.24	-0.147
MS-GARCH	1.934 (1.577)	0.568 (0.118)	50.66	-0.180
Japanese yen/U.S. dollar				
Binomial multifractal	-1.248 (2.160)	0.909 (0.155)	76.95	0.205
GARCH	5.311 (1.233)	0.488 (0.086)	99.15	-0.024
MS-GARCH	2.148 (1.776)	0.573 (0.108)	103.29	-0.067
British pound/U.S. dollar				
Binomial multifractal	0.330 (1.114)	0.792 (0.120)	27.35	0.250
GARCH	2.702 (0.760)	0.606 (0.085)	29.61	0.188
MS-GARCH	0.641 (1.021)	0.730 (0.105)	29.08	0.203
Canadian dollar/U.S. dollar				
Binomial multifractal	-0.038 (0.385)	1.179 (0.221)	1.6339	0.217
GARCH	0.676 (0.243)	0.707 (0.121)	1.6615	0.204
MS-GARCH	0.630 (0.270)	0.754 (0.140)	1.6719	0.199

This table gives out-of-sample forecasting results for the three models. The first two columns correspond to parameter estimates from the Mincer-Zarnowitz OLS regression $e_{t+1}^2 = \gamma_0 + \gamma_1 E_t(e_{t+1}^2) + u_t$. For an unbiased forecast we expect $\gamma_0 = 0$ and $\gamma_1 = 1$. Asymptotic standard errors in parentheses are corrected for heteroscedasticity and autocorrelation using the method of Newey and West (1987, 1994) and for parameter uncertainty using the method of West and McCracken (1998). MSE is the mean squared forecast error and R^2 is one less the MSE divided by the sum of squared demeaned squared returns in the out-of-sample period.

$\gamma_0 = 0$ and $\gamma_1 = 1$ in all cases at the 5% confidence level. In contrast, for the other models, each currency leads to a strong rejection of either one hypothesis (MS-GARCH) or both (GARCH) at the 5% confidence level. The reported MSE and R^2 further confirm that the multifractal provides the best 20-day forecasts for all currencies. The difference is particularly large in the case of the DEM and JPY.

The R^2 coefficient is 13.5% and 20.5% for DEM and JPY, respectively, with the multifractal, while negative values are produced by GARCH and MS-GARCH.¹⁴

Table 9 reports summary forecasting results and significance tests for horizons of 1, 5, 10, 20, and 50 days. Panel A shows the forecasting R^2 for each model. For the DEM and the JPY, the multifractal model is quite dominant at the 5-day horizon, and increasingly outperforms other models at longer horizons. For the GBP and the CAD, the multifractal only dominates at horizons of 20 days or more. Panel B analyzes the statistical significance of these results. At horizons of 50 days, the multifractal model outperforms the other models very significantly for the DEM, GBP, and JPY, and at a moderate or marginal significance level for CAD. The superior forecasts of the multifractal are also highly significant at horizons of 10 and 20 days for the DEM, and somewhat strong at the 20-day horizons for the JPY and GBP.

These results are quite impressive for the multifractal model. Although this is the first forecasting evaluation of MSM, and only the simplest binomial specification has been investigated, our process compares well with established models. In particular, GARCH(1,1) is often viewed as a standard benchmark that is very difficult to outperform in forecasting exercises [e.g., West and Cho (1995); Andersen and Bollerslev (1998); Hansen and Lunde (2001)]. Our results show that MSM matches or slightly outperforms GARCH and MS-GARCH at short horizons and substantially dominates these models at longer horizons.

4.3 Comparison with FIGARCH

The out-of-sample results suggest that MSM accurately captures the dependence structure of volatility at long horizons. Next we investigate whether a fractionally integrated volatility process also provides good long-range forecasts. We consider the FIGARCH (1, d , 0) specification of Baillie, Bollerslev, and Mikkelsen (1996).¹⁵ The return process is given by $x_t = h_t^{1/2} e_t$, where $\{e_t\}$ are i.i.d. Student innovations with unit variance and ν degrees of freedom. The conditional variance h_t satisfies

$$h_{t+1} = \omega + \beta(h_t - x_t^2) + [1 - (1 - L)^d]x_t^2,$$

where L denotes the lag operator and $d \in [0, 1]$ the long-memory parameter. FIGARCH is well defined, strictly stationary, and ergodic when $\omega \geq 0$ and $|\beta| < 1$. For every $d > 0$, the process is not covariance-stationary because the unconditional variance is infinite, as discussed in Baillie, Bollerslev, and Mikkelsen (1996).

We estimate FIGARCH by ML and report the corresponding results in Table 10, panel A.¹⁶ For every currency, FIGARCH has the lowest in-sample

¹⁴ The multifractal yields a higher R^2 for 20-day returns than for daily returns. This stems from the fact that our measure of 20-day volatility is a sum of daily squared returns $\sum_{s=t}^{t+19} x_s^2$. As in Andersen and Bollerslev (1998), reduced noise in the volatility measure leads to an increase in explanatory power.

¹⁵ In unreported work, FIGARCH(1, d , 0) was not rejected in favor of more general FIGARCH(p , d , q) specifications for any of the exchange rate series.

¹⁶ We note that for the JPY, the estimated value of $\omega = 0$ is on the boundary. For standard GARCH, this parameter determines the unconditional volatility level. Since in FIGARCH unconditional volatility is infinite, interpretation of this result is more problematic. We can view our reported estimates for the JPY as corresponding to an earlier specification with $\omega = 0$ introduced by Robinson (1991).

Table 9 Forecast summary, multiple horizons.

	Horizon (days)				
	1	5	10	20	50
Panel A. Restricted R^2					
Deutsche mark/U.S. dollar					
Binomial multifractal	0.041	0.124	0.160	0.135	0.038
GARCH	0.035	0.069	0.033	-0.147	-0.761
MS-GARCH	0.039	0.072	0.030	-0.180	-1.137
Japanese yen/U.S. dollar					
Binomial multifractal	0.053	0.113	0.142	0.205	0.213
GARCH	0.048	0.054	0.011	-0.024	-0.358
MS-GARCH	0.048	0.044	-0.009	-0.067	-0.569
British pound/U.S. dollar					
Binomial multifractal	0.057	0.165	0.235	0.250	0.273
GARCH	0.076	0.191	0.244	0.188	-0.026
MS-GARCH	0.072	0.165	0.238	0.203	0.038
Canadian dollar/U.S. dollar					
Binomial multifractal	0.051	0.172	0.221	0.217	0.111
GARCH	0.042	0.154	0.205	0.204	0.070
MS-GARCH	0.055	0.181	0.229	0.199	0.036
Panel B. MSE test vs. multifractal (p -value)					
Deutsche mark/U.S. dollar					
GARCH	0.307	0.040	0.009	0.001	0.000
MS-GARCH	0.314	0.004	0.000	0.000	0.000
Japanese yen/U.S. dollar					
GARCH	0.426	0.208	0.144	0.117	0.063
MS-GARCH	0.415	0.143	0.071	0.021	0.000
British pound/U.S. dollar					
GARCH	0.906	0.824	0.606	0.156	0.016
MS-GARCH	0.857	0.499	0.547	0.108	0.000
Canadian dollar/U.S. dollar					
GARCH	0.294	0.3590	0.410	0.447	0.292
MS-GARCH	0.597	0.603	0.565	0.380	0.065

This table summarizes out-of-sample forecasting results across multiple horizons. Panel A gives the restricted forecasting R^2 for each model and horizon. Panel B gives p -values from testing that the corresponding model has a lower out-of-sample forecasting MSE than the binomial multifractal. The tests are corrected for autocorrelation and heteroscedasticity using Newey and West (1987, 1994). A low p -value indicates that forecasts from the corresponding model would be rejected in favor of multifractal forecasts.

likelihood of all estimated models. In particular, the reported p -values indicate a difference in likelihood relative to MSM that is statistically significant at the 1% level for the JPY and at the 10% level for the other currencies. MSM thus outperforms FIGARCH in-sample.

Table 10 FIGARCH comparison.

Panel A. In-sample						
Parameter estimates				ln L	p -value vs. MSM	
	ω	β	d			
DEM	0.003 (0.001)	0.906 (0.028)	0.994 (0.045)	0.195 (0.011)	−5731.93	0.056
JPY	0 (−)	0.930 (0.006)	1.000 (−)	0.228 (0.008)	−5974.64	0.006
GBP	0.000 (0.000)	0.931 (0.005)	1.000 (−)	0.202 (0.008)	−5567.80	0.053
CAD	0.005 (0.001)	0.236 (0.034)	0.347 (0.028)	0.148 (0.011)	−105.96	0.062
Panel B. Out-of-sample						
Horizon (days)						
	1	5	10	20	50	
Restricted R^2						
DEM	0.022	0.065	0.080	0.028	−0.167	
JPY	0.042	0.009	−0.076	−0.153	−0.588	
GBP	0.074	0.183	0.231	0.167	−0.071	
CAD	0.030	0.152	0.232	0.246	0.142	
MSE Test vs. multifractal (p -value)						
DEM	0.125	0.050	0.056	0.046	0.014	
JPY	0.350	0.127	0.087	0.073	0.044	
GBP	0.882	0.716	0.462	0.102	0.011	
CAD	0.047	0.275	0.609	0.703	0.738	

This table gives FIGARCH estimation results and compares in- and out-of-sample results with the multifractal model. In Panel A, the first four columns give parameter estimates for FIGARCH(1, d , 0). In all cases, this specification could not be rejected in favor of more general FIGARCH(p , d , q). Asymptotic standard errors are in parentheses. For both JPY and GBP, the estimated value of d is on the boundary of $1 - \epsilon$, where $\epsilon = 10^{-5}$. For JPY, the estimated value of ω is on the boundary $\omega = 0$, corresponding to a long-memory volatility process suggested by Robinson (1991). We also report the value of the log-likelihood and a p -value from a test of whether the FIGARCH likelihood dominates MSM. The test corresponds to Vuong(1989) adjusted for heteroscedasticity and autocorrelation. In panel B we report out-of-sample forecasting R^2 statistics and their associated p -values, HAC adjusted. For all tests, a low p -value indicates that FIGARCH would be rejected in favor of MSM.

We report in Table 10, panel B, the out-of-sample results for FIGARCH. The restricted R^2 show that the multifractal outperforms FIGARCH at all horizons. MSM dominates at short horizons (1 and 5 days) for the DEM, JPY, and CAD, and at long horizons (10, 20, and 50 days) for the DEM, JPY, and GBP. Despite its long-memory structure, FIGARCH performs especially poorly at 50-day horizons. The corresponding R^2 are negative for the DEM, JPY, and GBP. The p -values of the

MSEs confirm the statistical significance of these results. At the 10% confidence level, 9 of the 16 MSM forecasts significantly outperform FIGARCH, while none of the FIGARCH forecasts significantly improve on MSM. Overall the results of Table 10 show that MSM dominates FIGARCH both in- and out-of-sample.

These results indicate that long memory alone does not explain the excellent long-run forecasting performance of the multifractal model. As suggested in subsection 1.2, fractional integration as introduced in FIGARCH tends to produce smooth patterns in volatility. In contrast, our approach generates long cycles with a switching mechanism that also gives abrupt volatility changes. Generating long memory with multiplicative shocks of heterogeneous frequencies is not only intuitively appealing, it also generates the best long-run forecasts.

5 CONCLUSION

This article proposes an expanded role for regime switching in volatility modeling. Traditional approaches, such as Markov switching ARCH [Cai (1994), Hamilton and Susmel (1994)] and GARCH [Gray (1996), Klaassen (2002)], consider separately three categories of volatility dynamics. High-frequency variation is captured by a thick-tailed conditional distribution of returns, midrange frequencies by smooth ARCH or GARCH components, and only very low frequencies are modeled with regime switching. We suggest an alternative approach based on regime switching at all frequencies. The model is tightly parameterized in spite of a high-dimensional state space. Using four long series of daily exchange rates, we find that MSM matches or dominates the performance of previous models across a range of in-sample and out-of-sample measures. Thus the primary contribution of the article is to show that regime switching can have a much broader scope than previously envisioned. In particular, our pure regime switching model provides a viable alternative to approaches that combine regime switching, linear volatility dynamics, and flexible tail distributions.

Researchers often focus on applications of immediate practical value when assessing statistical models. We have similarly shown that MSM does well by several standard measures of performance. From a theoretical perspective, good econometric descriptions are also useful as the first step in explanation. Whereas the standard approach separately addresses three distinct statistical phenomena, MSM offers a unified modeling strategy. As a result, this article invites the theorist to explore the economic mechanisms causing self-similar regime switching in financial volatility.

Returning to the more immediate contributions of our work, this article develops the first comprehensive econometric toolkit to estimate and test multifractal processes. We develop a ML estimator and show that it works well in Monte Carlo simulations. The application to exchange rates indicates that the likelihood function increases with the number of volatility components. This finding is important because it confirms substantial heterogeneity in volatility persistence, which is one of the primary motivations of the multifractal approach.

Consistent with intuition, the spacing of volatility components across frequencies becomes tighter and the contribution of individual volatility components becomes smaller as the number of components increases.

We also compare the multifractal model with GARCH, MS-GARCH and FIGARCH, which are chosen because of demonstrated good performance with exchange rate data under a variety of metrics. Like MSM, these models conveniently permit ML estimation and analytical forecasting. In-sample, the likelihood is significantly higher for the multifractal than for GARCH and FIGARCH, even though these processes have the same number of parameters. We use a BIC criterion to compare goodness-of-fit of MS-GARCH and MSM, and find no significant difference between the two models. Out-of-sample evidence further validates the multifractal approach. MSM matches or outperforms GARCH, MS-GARCH, and FIGARCH at short horizons and substantially dominates these at longer horizons. The multifractal model thus offers a promising new approach to volatility modeling and is well deserving of further empirical and theoretical investigation.

APPENDIX

A.1 Proof of Proposition 1

Consider a sequence of processes with fixed parameter vector $\psi = (m_0, \sigma, b, \gamma^*)$. Note in particular that $\gamma_{\bar{k}} = \gamma^*$ for all \bar{k} . For any integer $n \geq 0$ and real $q \in [0, \infty)$, it is convenient to define $K_q(n) = \mathbb{E}(|x_t|^q |x_{t+n}|^q) / [\mathbb{E}(|x_t|^{2q})]$ and $c_q = [\mathbb{E}(|\varepsilon_t|^q)]^2 / [\mathbb{E}(|\varepsilon_t|^{2q})]$. Multipliers in different stages of the cascade are statistically independent. The definition of returns, $x_t = \sigma(M_{1,t}M_{2,t} \dots M_{\bar{k},t})^{1/2} \varepsilon_t$ and $x_{t+n} = \sigma(M_{1,t+n}M_{2,t+n} \dots M_{\bar{k},t+n})^{1/2} \varepsilon_{t+n}$, implies

$$K_q(n) = c_q [\mathbb{E}(M^q)]^{-\bar{k}} \prod_{k=1}^{\bar{k}} \mathbb{E}(M_{k,t}^{q/2} M_{k,t+n}^{q/2}).$$

Note that $\mathbb{E}(M_{k,t}^{q/2} M_{k,t+n}^{q/2}) = \mathbb{E}(M^q)(1 - \gamma_k)^n + [\mathbb{E}(M^{q/2})]^2 [1 - (1 - \gamma_k)^n]$, or equivalently

$$\mathbb{E}(M_{k,t}^{q/2} M_{k,t+n}^{q/2}) = [\mathbb{E}(M^{q/2})]^2 [1 + a_q(1 - \gamma_k)^n],$$

where $a_q = \mathbb{E}(M^q) / [\mathbb{E}(M^{q/2})]^2 - 1$. Since $1 - \gamma_k = (1 - \gamma^*)^{b^k - k}$ and $\gamma_{\bar{k}} = \gamma^*$, we obtain

$$\ln \frac{K_q(n)}{c_q} = \sum_{k=1}^{\bar{k}} \ln \frac{1 + a_q(1 - \gamma^*)^{nb^k - k}}{1 + a_q} \quad (6)$$

As k increases from 1 to \bar{k} , the expression $(1 - \gamma^*)^{nb^k - k}$ declines from $(1 - \gamma^*)^{nb^{1-k}}$ to $(1 - \gamma^*)^n$. The maximum $(1 - \gamma^*)^{nb^{1-k}}$ is close to one and the minimum $(1 - \gamma^*)^n$ is close to zero when $b^{\bar{k}}/n$ and n are large. Intermediate values are observed when $(1 - \gamma^*)^{nb^k - k} \approx 1 - \gamma^*$, or equivalently $k \approx \log_b(b^{\bar{k}}/n)$. Let $i(n)$ denote the unique integer such that $i(n) \leq \log_b(b^{\bar{k}}/n) < i(n) + 1$. We anticipate that

$$\ln \frac{K_q(n)}{c_q} \approx \sum_{k=i(n)+1}^{\bar{k}} \ln \frac{1}{1 + a_q} = -[\bar{k} - i(n)] \ln(1 + a_q),$$

and thus $\ln K_q(n) \approx -(\log_b n) \ln(1 + a_q) = -\delta(q) \ln n$.

To formalize this intuition, consider the interval $I_{\bar{k}} = \{n: \alpha_1 \log_b(b^{\bar{k}}) \leq \log_b n \leq \alpha_2 \log_b(b^{\bar{k}})\}$. Note that $\log_b(b^{\bar{k}}/n) \geq (1 - \alpha_2) \log_b(b^{\bar{k}})$ for all $n \in I_{\bar{k}}$. We henceforth assume that \bar{k} is sufficiently large so that $i(n) \geq b \forall n \in I_{\bar{k}}$. Consider an arbitrary sequence of strictly positive integers $j(n)$ monotonically diverging to $+\infty$. The precise definition of $j(n)$ is temporarily postponed. Let

$$u_n = j(n) \ln(1 + a_q) + \sum_{k=i(n)-j(n)+1}^{i(n)+j(n)} \ln \frac{1 + a_q(1 - \gamma^*)^{nb^{k-\bar{k}}}}{1 + a_q}.$$

By Equation (6), $\ln[K_q(n)/c_q]$ can be decomposed into four components:

$$\begin{aligned} \ln \frac{K_q(n)}{c_q} = & -[\bar{k} - i(n)] \ln(1 + a_q) + \sum_{k=1}^{i(n)-j(n)} \ln \frac{1 + a_q(1 - \gamma^*)^{nb^{k-\bar{k}}}}{1 + a_q} \\ & + u_n + \sum_{k=i(n)+j(n)+1}^{\bar{k}} \ln[1 + a_q(1 - \gamma^*)^{nb^{k-\bar{k}}}]. \end{aligned} \quad (7)$$

We successively examine each component on the right-hand side.

- The first component is between $-\delta(q)(\ln n + \ln b)$ and $-\delta(q) \ln n$.
- The second component contains terms $(1 - \gamma^*)^{nb^{k-\bar{k}}}$ that are bounded below by $(1 - \gamma^*)^{nb^{i(n)-j(n)-\bar{k}}}$. The definition of $i(n)$ implies $nb^{i(n)-\bar{k}} \leq 1$ and thus

$$\left| \sum_{k=1}^{i(n)-j(n)} \ln \frac{1 + a_q(1 - \gamma^*)^{nb^{k-\bar{k}}}}{1 + a_q} \right| \leq i(n) \ln \frac{1 + a_q}{1 + a_q(1 - \gamma^*)^{b^{-j(n)}}}.$$

By standard concavity arguments, the second component of Equation (7) is therefore bounded by $i(n)b^{-j(n)}a_q |\ln(1 - \gamma^*)|$.

- The third component, u_n , contains terms $1 + a_q(1 - \gamma^*)^{nb^{k-\bar{k}}}$ that are between 1 and $1 + a_q$. Hence $|u_n| \leq j(n) \ln(1 + a_q) \leq a_q j(n)$.
- The fourth component is positive and bounded above by

$$a_q \sum_{k=i(n)+j(n)+1}^{\bar{k}} (1 - \gamma^*)^{nb^{k-\bar{k}}} \leq a_q \sum_{k=0}^{\infty} (1 - \gamma^*)^{b^k nb^{i(n)+j(n)+1-\bar{k}}}.$$

We check that $nb^{i(n)+j(n)+1-\bar{k}} \geq 1$ and $b^k \geq k(b-1)$. The fourth component is therefore bounded above by $a_q \sum_{k=0}^{\infty} (1 - \gamma^*)^{k(b-1)} = a_q / [1 - (1 - \gamma^*)^{b-1}]$.

This establishes that

$$\left| \frac{\ln K_q(n)}{\ln n^{-\delta(q)}} - 1 \right| \leq \frac{c_q^* + a_q j(n) + a_q i(n) b^{-j(n)} |\ln(1 - \gamma^*)|}{\delta(q) \ln n},$$

where $c_q^* = \delta(q) \ln b + |\ln c_q| + a_q/[1 - (1 - \gamma^*)^{(b-1)}]$. We now choose a sequence $j(n)$ such that the right-hand side of the inequality converges to zero. More specifically, consider the unique integer such that¹⁷ $j(n) \leq 2 \log_b i(n) < j(n) + 1$. It is easy to check that $i(n)b^{-j(n)} = b^{\log_b i(n) - j(n)} \leq 1$ and $j(n) \leq 2 \log_b (\log_b b^{\bar{k}}) = 2 \log_b \bar{k}$. For all $n \in I_{\bar{k}}$, the quantity $|(\ln K_q(n)/\ln n^{-\delta(q)}) - 1|$ is therefore bounded above by

$$\eta_{\bar{k}} = \frac{1}{\bar{k} \delta(q) \alpha_1 \ln b} [2a_q \log_b \bar{k} + c_q^* + a_q |\ln(1 - \gamma^*)|], \quad (8)$$

which is independent of n . We infer that $\sup_{n \in I_{\bar{k}}} |(\ln K_q(n)/\ln n^{-\delta(q)}) - 1| \rightarrow 0$ as $\bar{k} \rightarrow +\infty$.

Finally, it is easy to show that the autocorrelation $\rho_q(n)$ satisfies

$$1 \leq \frac{K_q(n)}{\rho_q(n)} = \frac{1 - c_q(1 + a_q)^{-\bar{k}}}{1 - c_q(1 + a_q)^{-\bar{k}}/K_q(n)} \leq \frac{1}{1 - c_q(1 + a_q)^{-\bar{k}}/K_q(n)}. \quad (9)$$

Equation (8) implies that for all $n \in I_{\bar{k}}$, $\log_b K_q(n) \geq -\delta(q)(1 + \eta_{\bar{k}})\alpha_2 \bar{k}$, and thus

$$\log_b [K_q(n)/(1 + a_q)^{-\bar{k}}] \geq \bar{k} \delta(q)(1 - \alpha_2 - \alpha_2 \eta_{\bar{k}}/\bar{k}). \quad (10)$$

Combining Equations (9) and (10), we conclude that $\sup_{n \in I_{\bar{k}}} |\ln(K_q(n)/\rho_q(n))| \rightarrow 0$ and thus that the proposition holds.

A.2 HAC-Adjusted Vuong Test

We consider the probability space $(\Omega, \mathcal{F}, \mathbb{P}^0)$ and a stochastic process $\{x_t\}_{t=-\infty}^{+\infty}$. Each X_t is a random variable taking values on the real line. For every t , it is convenient to consider the vector of past values $X_{t-1} = \{x_s\}_{s=-\infty}^{t-1}$. The econometrician directly observes a finite number of realizations of x_t , but ignores the true data-generating process. She instead considers two competing families of models specified by their conditional densities $\mathcal{M}_f = \{f(x_t | X_{t-1}, \theta); \theta \in \Theta\}$ and $\mathcal{M}_g = \{g(x_t | X_{t-1}, \gamma); \gamma \in \Gamma\}$. These families may or may not contain the true data-generating process. The pseudo-true value θ^* specifies the model in \mathcal{M}_f with the optimal Kullback-Leibler information criterion:

$$\theta^* = \arg \max_{\theta \in \Theta} \mathbb{E}^0 [\ln f(x_t | X_{t-1}, \theta)].$$

The pseudo-true-value γ^* is similarly defined.

Consider the log-likelihood functions:

$$L_T^f(\theta) \equiv \sum_{t=1}^T \ln f(x_t | X_{t-1}, \theta), \quad L_T^g(\gamma) \equiv \sum_{t=1}^T \ln g(x_t | X_{t-1}, \gamma).$$

By definition, the ML estimators $\hat{\theta}_T$ and $\hat{\gamma}_T$ maximize the functions $L_T^f(\theta)$ and $L_T^g(\gamma)$. The corresponding first order conditions are

$$\frac{\partial L_T^f}{\partial \theta}(\hat{\theta}_T) = 0, \quad \frac{\partial L_T^g}{\partial \gamma}(\hat{\gamma}_T) = 0. \quad (11)$$

¹⁷ We check that when \bar{k} is large enough, $1 \leq j(n) \leq i(n)$ and $j(n) + i(n) \leq \bar{k}$ for all $n \in I_{\bar{k}}$.

We now examine the likelihood ratio

$$LR_T(\hat{\theta}_T, \hat{\gamma}_T) = L_T^f(\hat{\theta}_T) - L_T^g(\hat{\gamma}_T) = \sum_{t=1}^T \ln \frac{f(x_t | X_{t-1}, \hat{\theta}_T)}{g(x_t | X_{t-1}, \hat{\gamma}_T)}.$$

By Equation (11), a second-order expansion of LR_T implies that $(1/\sqrt{T})LR_T(\hat{\theta}_T, \hat{\gamma}_T) = 1/\sqrt{T}LR_T(\theta^*, \gamma^*) + o_p(1)$, and thus

$$\frac{1}{\sqrt{T}}LR_T(\hat{\theta}_T, \hat{\gamma}_T) = \frac{1}{\sqrt{T}} \sum_{t=1}^T \ln \frac{f(x_t | X_{t-1}, \theta^*)}{g(x_t | X_{t-1}, \gamma^*)} + o_p(1).$$

Let $a_t = \ln[f(x_t | X_{t-1}, \theta^*)/g(x_t | X_{t-1}, \gamma^*)]$ and $\hat{a}_t = \ln[f(x_t | X_{t-1}, \hat{\theta}_T)/g(x_t | X_{t-1}, \hat{\gamma}_T)]$.

When the observations x_t are i.i.d., the addends a_t are also i.i.d. If the models f and g have equal Kullback-Leibler information criterion, the Central Limit Theorem implies $(1/\sqrt{T})LR_T(\hat{\theta}_T, \hat{\gamma}_T) \xrightarrow{d} \mathcal{N}(0, \sigma_*^2)$, where $\sigma_*^2 = \text{var}(a_t)$. The variance is consistently estimated by the sample variance of $\{\hat{a}_t\}$.

In the non-i.i.d. case, we need to adjust for the correlation in the addends a_t . Let $\sigma_T^2 = 1/T \sum_{s=1}^T \sum_{t=1}^T \mathbb{E}(a_s a_t)$. We know that $1/\sqrt{T}LR_T(\hat{\theta}_T, \hat{\gamma}_T) = \sigma_T Z + o_p(1)$, where Z is a standard Gaussian. Following Newey and West (1987), we estimate σ_T by

$$\hat{\sigma}_T^2 = \hat{\Omega}_0 + 2 \sum_{j=1}^{m_T} w(j, m) \hat{\Omega}_j,$$

where $\hat{\Omega}_j = \sum_{t=j+1}^T \hat{a}_t \hat{a}_{t-j} / T$ denotes the sample covariance of $\{\hat{a}_t\}$, and $w(j, m) = 1 - j/(m+1)$ is the Bartlett weight. We choose m_T using the automatic lag selection method of Newey and West (1994).

Received October 25, 2002; revised March 4, 2003; accepted October 29, 2003

REFERENCES

- Akgiray, V. (1989). "Conditional Heteroskedasticity in Time Series of Stock Returns: Evidence and Forecasts." *Journal of Business* 62, 55–80.
- Albert, J., and S. Chib. (1993). "Bayes Inference via Gibbs Sampling of Autoregressive Time Series Subject to Markov Means and Variance Shifts." *Journal of Business and Economic Statistics* 11, 1–15.
- Andersen, T., and T. Bollerslev. (1998). "Answering the Skeptics: Yes, Standard Volatility Models Do Provide Accurate Forecasts." *International Economic Review* 39, 885–905.
- Andersen, T., T. Bollerslev, F. Diebold, and P. Labys. (2001). "The Distribution of Realized Exchange Rate Volatility." *Journal of the American Statistical Association* 96, 42–55.
- Andersen, T., T. Bollerslev, and N. Meddahi. (2002). "Analytic Evaluation of Volatility Forecasts." Working Paper 2002s-90. CIRANO; forthcoming in *International Economic Review*.
- Andrews, D. (1991). "Heteroskedasticity and Autocorrelation Consistent Covariance Matrix Estimation." *Econometrica* 59, 817–854.

- Andrews, D., and C. Monahan. (1992). "An Improved Heteroskedasticity and Autocorrelation Consistent Covariance Matrix Estimator." *Econometrica* 60, 953–966.
- Baillie, R., T. Bollerslev, and H. O. Mikkelsen. (1996). "Fractionally Integrated Generalized Autoregressive Conditional Heteroscedasticity." *Journal of Econometrics* 74, 3–30.
- Barndorff-Nielsen, O., and N. Shephard. (2003). "Realized Power Variation and Stochastic Volatility Models." *Bernoulli* 9, 243–265.
- Baum, L., T. Petrie, G. Soules, and N. Weiss. (1980). "A Maximization Technique Occurring in the Statistical Analysis of Probabilistic Functions of Markov Chains." *Annals of Mathematical Statistics* 41, 164–171.
- Bhattacharya, R., V. Gupta, and E. Waymire. (1983). "The Hurst Effect Under Trends." *Journal of Applied Probability* 20, 649–662.
- Bollen, N., S. Gray, and R. Whaley. (2000). "Regime Switching in Foreign Exchange Rates: Evidence from Currency Option Prices." *Journal of Econometrics* 94, 239–276.
- Bollerslev, T. (1987). "A Conditional Heteroskedastic Time Series Model for Speculative Prices and Rates of Return." *Review of Economics and Statistics* 69, 542–547.
- Cai, J. (1994). "A Markov Model of Switching-Regime ARCH." *Journal of Business and Economic Statistics* 12, 309–316.
- Calvet, L., and A. Fisher. (2001). "Forecasting Multifractal Volatility." *Journal of Econometrics* 105, 27–58.
- Calvet, L., and A. Fisher. (2002a). "Multifractality in Asset Returns: Theory and Evidence." *Review of Economics and Statistics* 84, 381–406.
- Calvet, L., and A. Fisher. (2002b). "Regime Switching and the Estimation of Multifractal Processes." Working Paper 02-064, New York University.
- Calvet, L., A. Fisher, and B. Mandelbrot. (1997). Cowles Foundation Discussion Papers no. 1164–1166, Yale University. Available at <http://www.ssrn.com>.
- Dacorogna, M., U. Müller, R. Nagler, R. Olsen, and O. Pictet. (1993). "A Geographical Model for the Daily and Weekly Seasonal Volatility in the Foreign Exchange Market." *Journal of International Money and Finance* 12, 413–438.
- den Haan, W., and A. Levin. (1997). "A Practitioner's Guide to Robust Covariance Matrix Estimation." In G. S. Maddala and C. R. Rao (eds.), *Handbook of Statistics: Robust Inference*, vol. 15. New York: Elsevier.
- Diebold, F., and A. Inoue. (2001). "Long Memory and Regime Switching." *Journal of Econometrics* 105, 27–58.
- Diebold, F., J.-H. Lee, and G. Weinbach. (1994). "Regime Switching with Time-Varying Transition Probabilities." In Colin Hargreaves (ed.), *Nonstationary Time Series Analysis and Cointegration*. New York: Oxford University Press.
- Ding, Z., C. Granger, and R. Engle. (1993). "A Long Memory Property of Stock Market Returns and a New Model." *Journal of Empirical Finance* 1, 83–106.
- Durland, M., and T. McCurdy. (1994). "Duration-Dependent Transitions in a Markov Model of US GNP Growth." *Journal of Business and Economic Statistics* 12, 279–288.
- Engle, R. (1982). "Autoregressive Conditional Heteroscedasticity with Estimates of the Variance of United Kingdom Inflation." *Econometrica* 50, 987–1007.
- Filardo, A. (1994). "Business Cycle Phases and Their Transitional Dynamics." *Journal of Business and Economic Statistics* 12, 299–308.
- Garcia, R. (1998). "Asymptotic Null Distribution of the Likelihood Ratio Test in Markov Switching Models." *International Economic Review* 39, 763–788.

- Garcia, R., and P. Perron. (1996). "An Analysis of the Real Interest Rate Under Regime-Shifts." *Review of Economics and Statistics* 78, 111–125.
- Ghysels, E., A. Harvey, and E. Renault. (1996). "Stochastic Volatility." In G. S. Maddala and C. R. Rao (eds.), *Handbook of Statistics*, vol. 14. Amsterdam: North-Holland.
- Gouriéroux, C., and J. Jasiak. (2002). "Nonlinear Autocorrelograms: An Application to Inter-Trade Durations." *Journal of Time Series Analysis* 23, 127–154.
- Gray, S. (1996). "Modeling the Conditional Distribution of Interest Rates as a Regime-Switching Process." *Journal of Financial Economics* 42, 27–62.
- Hamilton, J. (1988). "Rational Expectations Econometric Analysis of Changes in Regimes: An Investigation of the Term Structure of Interest Rates." *Journal of Economic Dynamics and Control* 12, 385–423.
- Hamilton, J. (1989). "A New Approach to the Economic Analysis of Nonstationary Time Series and the Business Cycle." *Econometrica* 57, 357–384.
- Hamilton, J. (1990). "Analysis of Time Series Subject to Change in Regime." *Journal of Econometrics* 45, 39–70.
- Hamilton, J. (1994). *Time Series Analysis*. Princeton, NJ: Princeton University Press.
- Hamilton, J., and G. Lin. (1996). "Stock Market Volatility and the Business Cycle." *Journal of Applied Econometrics* 11, 573–593.
- Hamilton, J., and G. Pérez-Quirós. (1996). "What Do the Leading Indicators Lead?" *Journal of Business* 69, 27–49.
- Hamilton, J., and B. Raj. (2002). "New Directions in Business Cycle Research and Financial Analysis." *Empirical Economics* 27, 149–162.
- Hamilton, J., and R. Susmel. (1994). "Autoregressive Conditional Heteroskedasticity and Changes in Regime." *Journal of Econometrics* 64, 307–333.
- Hansen, B. (1992). "The Likelihood Ratio Test Under Non-Standard Conditions: Testing the Markov switching Model of GNP." *Journal of Applied Econometrics* 7, 561–582.
- Hansen, P., and A. Lunde. (2001). "A Forecast Comparison of Volatility Models: Does Anything Beat a GARCH(1,1)?" Working paper, Brown University.
- Hidalgo, J., and P. Robinson. (1996). "Testing for Structural Change in a Long-Memory Environment." *Journal of Econometrics* 70, 159–174.
- Kim, C.-J. (1994). "Dynamic Linear Models with Markov switching." *Journal of Econometrics* 60, 1–22.
- Kim, C.-J., and C. Nelson. (1999). *State-Space Models with Regime Switching*. Cambridge, MA: MIT Press.
- Klaassen, F. (2002). "Improving GARCH Volatility Forecasts with Regime Switching GARCH." *Empirical Economics* 27, 363–394.
- LeBaron, B. (2001). "Stochastic Volatility as a Simple Generator of Apparent Financial Power Laws and Long Memory." *Quantitative Finance* 1, 621–631.
- Lindgren, G. (1978). "Markov Regime Models for Mixed Distributions and Switching Regressions." *Scandinavian Journal of Statistics* 5, 81–91.
- Maheu, J., and T. McCurdy. (2000). "Volatility Dynamics Under Duration-Dependent Mixing." *Journal of Empirical Finance* 7, 345–372.
- Newey, W., and K. West. (1987). "A Simple, Positive Semi-Definite, Heteroskedasticity and Autocorrelation Consistent Covariance Matrix." *Econometrica* 55, 703–708.
- Newey, W., and K. West. (1994). "Automatic Lag Selection in Covariance Matrix Estimation." *Review of Economic Studies* 61, 631–654.
- Pagan, A., and W. Schwert. (1990). "Alternative Models for Conditional Stock Volatility." *Journal of Econometrics* 45, 267–290.

- Pérez-Quirós, G., and A. Timmerman. (2000). "Firm Size and Cyclical Variations in Stock Returns." *Journal of Finance* 55, 1229–1262.
- Rivers, D., and Q. Vuong. (2002). "Model Selection Tests for Nonlinear Dynamic Models." *Econometrics Journal* 5, 1–39.
- Robinson, P. (1991). "Testing for Strong Serial Correlation and Dynamic Conditional Heteroskedasticity in Multiple Regression." *Journal of Econometrics* 47, 67–84.
- Shephard, N. (1994). "Partial Non-Gaussian State Space." *Biometrika* 81, 115–131.
- Vuong, Q. (1989). "Likelihood Ratio Tests for Model Selection and Non-Nested Hypotheses." *Econometrica* 57, 307–333.
- West, K., and D. Cho. (1995). "The Predictive Ability of Several Models of Exchange Rate Volatility." *Journal of Econometrics* 69, 367–391.
- West, K., and M. McCracken. (1998). "Regression-Based Tests of Predictive Ability." *International Economic Review* 39, 817–840.